

# Data Science for Economists

## Lecture 2: R language basics

---

Drew Van Kuiken

University of North Carolina | ECON 370

# Table of contents

1. Prologue
2. Introduction
3. Types of Data
4. Data Structures
5. Getting Help

\* Slides adapted from Grant McDermott's EC 607 at University of Oregon.

# Prologue

---

# Agenda

Today and the next lecture are going to be very hands on.

- I'll have slides as per usual, but we're going to spend a lot of time live coding together.

This is deliberate.

- I want you to get comfortable typing R commands yourself — and navigating the RStudio IDE — without resorting to copy+paste.
- Slightly more painful in the beginning, but much better payoff in the long-run.

# Introduction

---

(Some important R concepts)

# Basic Arithmetic

R is a powerful calculator and recognizes all of the standard arithmetic operators:

```
1+2 ## Addition
```

```
## [1] 3
```

```
6-7 ## Subtraction
```

```
## [1] -1
```

```
5/2 ## Division
```

```
## [1] 2.5
```

```
2^3 ## Exponentiation
```

```
## [1] 8
```

```
2+4*1^3 ## Standard order of precedence (`*` before `+`, etc.)
```

```
## [1] 6
```

# Basic Arithmetic

When possible, do operators in vectors in `R`.

```
first_vec = 1:5  #store our first vector  
first_vec      #show vector
```

```
## [1] 1 2 3 4 5
```

```
first_vec + 0.5  #add 1/2 to each element
```

```
## [1] 1.5 2.5 3.5 4.5 5.5
```

```
first_vec + 6:10 #add another vector to first_vector
```

```
## [1]  7  9 11 13 15
```

```
first_vec + 6:9  #oops!
```

```
## Warning in first_vec + 6:9: longer object length is not a multiple of shorter  
## object length
```

```
## [1]  7  9 11 13 11
```

# Basic Arithmetic (cont.)

We can also invoke modulo operators (quotient & remainder).

- Very useful when dealing with time, for example.

```
100 %/% 60 ## How many whole hours in 100 minutes?
```

```
## [1] 1
```

```
120 %/% 60 ## How many whole hours in 120 minutes?
```

```
## [1] 2
```

```
100 %% 60 ## How many minutes are left over from dividing 100 by 60?
```

```
## [1] 40
```

```
120 %% 60 ## How many minutes are left over from dividing 120 by 60?
```

```
## [1] 0
```



# Data Types in R

---

# Data Types

R has 6 basic data types:

1. Character
2. Numeric
3. Integer
4. Logical
5. Complex
6. Raw (we will mostly ignore this type)

# Data Types (cont.)

## Character

- The character type can be described as "text."
  - It is known as a "string" in other programming languages.

```
my_name      = "Drew Van Kuiken"  
first_name   = "Drew"  
last_name    = "Van Kuiken"      #good style to line up to equals signs (i often don't p  
class(my_name)
```

```
## [1] "character"
```

```
is.character(my_name)
```

```
## [1] TRUE
```

# Data Types (cont.)

## Character: Counting Characters

```
length(my_name)
```

```
## [1] 1
```

```
length(c(first_name,last_name))
```

```
## [1] 2
```

```
nchar(my_name)
```

```
## [1] 15
```

```
nchar(c(first_name,last_name))
```

```
## [1] 4 10
```

# Data Types (cont.)

## Character: Combining Characters

```
also_my_name = 'Drew Van Kuiken'  
my_name = also_my_name
```

```
## [1] TRUE
```

```
my_name_in_dutch = 'Drew van Kuiken'  
my_name = my_name_in_dutch
```

```
## [1] FALSE
```

```
paste(first_name,last_name)
```

```
## [1] "Drew Van Kuiken"
```

```
paste(first_name,last_name,sep="-")
```

```
## [1] "Drew-Van Kuiken"
```

```
paste0(first_name,last_name)
```

# Data Types (cont.)

## Character: An Aside

- The following code will look *very* similar to a user when examined in the console.

```
my_name
```

```
## [1] "Drew Van Kuiken"
```

```
print(my_name)
```

```
## [1] "Drew Van Kuiken"
```

- However, what's happening is ***very different***.
  - The first thing returns the *object* `my_name`
  - The second only prints the *contents* of the object `my_name`
- Imagine I have something written on a piece of paper. The first line is if I handed you the piece of paper and you read it. The second line is if I told you aloud what was written on the paper.
  - This difference will matter a lot when it comes to functions.

# Data Types (cont.)

## Character

```
first_name + " " + last_name
```

```
## Error in first_name + " ": non-numeric argument to binary operator
```

```
toupper(paste(first_name, last_name))
```

```
## [1] "DREW VAN KUIKEN"
```

```
tolower(paste(first_name, last_name))
```

```
## [1] "drew van kuiken"
```

# Data Types (cont.)

## Character: Special Characters

- Special characters "not native" to English can still be used but must be encoded so that `R` handles them correctly.
- Encoding guarantees that the computer will handle the special characters correctly.
- We will not spend much time on this; however, to read more about encoding text, [see this post](#).



# Data Types (cont.)

## Numeric

Numeric data are "numbers"

- In a lot of programming languages, there is this distinction between floats and integers. This is also true in R but to a much smaller degree.
- Unless explicitly stated as an integer, all "numbers" are numeric data.

```
my_age      = 29
my_height   = 6 + 4/12 #+1/12
c(my_age,my_height)
```

```
## [1] 29.000000 6.333333
```

```
class(my_age)
```

```
## [1] "numeric"
```

```
class(my_height)
```

```
## [1] "numeric"
```

# Data Types (cont.)

## Special Numeric Data

- There are a handful of special numeric values including Inf, -Inf and NaN
- Be careful when working with these

```
1/0
```

```
## [1] Inf
```

```
log(0)
```

```
## [1] -Inf
```

```
log(-1)
```

```
## Warning in log(-1): NaNs produced
```

```
## [1] NaN
```

# Data Types (cont.)

## Integers

"Whole numbers"

- Integers in R are distinguished from numeric data by having an L after the number part.
- The distinction between numeric and integers is not that important here in R.
  - We will mostly ignore the distinction.

```
also_my_age = 29L  
class(my_age)
```

```
## [1] "numeric"
```

```
class(also_my_age)
```

```
## [1] "integer"
```

# Data Types (cont.)

## Logical

Logical data are either `TRUE` or `FALSE`.

- `T` and `F` are equivalent.

```
R_is_fun  = TRUE
R_is_hard = FALSE
true      = T
false     = F
```

```
R_is_fun = true
```

```
## [1] TRUE
```

```
R_is_hard = false
```

```
## [1] TRUE
```

# Data Types (cont.)

## Logical: Operations

The most common use of logical data is for testing conditions and control flow.

- `>`, `<`, `>=`, and `<=` test for greater/less than and/or equal to.
- `&` and `|` are the logical operators for "and" and "or" respectively.
  - Order of operations: `&` are evaluated before `|`.
- `!` negates a logical: `!TRUE` becomes `FALSE` and vice versa,
- To test if equal, use two equal signs `==`. Not equal `!=`.

# Data Types (cont.)

## Logical: Operations

```
2 > 1
```

```
## [1] TRUE
```

```
1 > 2 & 1 > 1/2
```

```
## [1] FALSE
```

```
1 > 2 | 1 > 1/2
```

```
## [1] TRUE
```

```
2 > 1 | 1 > 1/2
```

```
## [1] TRUE
```

# Data Types (cont.)

## Logical: Operations

```
1 > 1/2 & 1 > 2 | 3 > 2    #order of operations are important
```

```
## [1] TRUE
```

```
1 > 2 & (1 > 1/2 | 3 > 2) #use parentheses when in doubt
```

```
## [1] FALSE
```

```
0.5 == 1/2
```

```
## [1] TRUE
```

```
3 != 2    #that is ! followed by a =
```

```
## [1] TRUE
```

# Data Types (cont.)

## Logical: Operations

`%in%` tests for containment

```
R_is_fun
```

```
## [1] TRUE
```

```
!R_is_fun
```

```
## [1] FALSE
```

```
2 %in% c(1,2,3,4)
```

```
## [1] TRUE
```

```
!(5 %in% c(1,2,3,4))
```

```
## [1] TRUE
```



# Aside: Binary operators

Operators like `+`, `-`, `/`, `%%`, `%in%`, etc are special functions in R that are binary operators. Notice that above to do "not in," it was a bit clunky:

```
!(5 %in% c(1,2,3,4))
```

You can define a new binary operator like so:

```
"%!in%" = function(a,b){!(a %in% b)}
```

```
5 %!in% c(1,2,3,4)
```

```
## [1] TRUE
```

# Data Types (cont.)

## Special Logicals

- `NA` is a special type of logical data
- Cannot be compared the same way.
- `NA` is different than `NaN`

```
NA == TRUE
```

```
## [1] NA
```

```
is.na(NA)
```

```
## [1] TRUE
```

```
1 == NaN
```

```
## [1] NA
```

```
is.nan(NA)
```

```
## [1] FALSE
```

# Data Types (cont.)

## Logical: Operations

If "mathematical operations" are performed on logicals, they will be coerced into 1s and 0s.

```
as.numeric(c(TRUE, FALSE))
```

```
## [1] 1 0
```

```
TRUE + FALSE
```

```
## [1] 1
```

```
sum(c(TRUE, FALSE, TRUE))
```

```
## [1] 2
```

```
mean(c(TRUE, FALSE, TRUE))
```

```
## [1] 0.6666667
```

# Data Types (cont.)

## Logical: Comparing Floats

Must be careful when testing for equality of floating point numbers ("decimals")

```
0.1 + 0.2 == 0.3
```

```
## [1] FALSE
```

What happened?

- Finite precision: floating point numbers aren't exact due to how numbers are stored in a computer.

```
all.equal(0.1+0.2,0.3)
```

```
## [1] TRUE
```

# Data Types (cont.)

## Logical

You can read more about logical operators and types [here](#) and [here](#).

# Data Types (cont.)

## Complex

Complex data are data that have complex numbers.

- We probably won't use these much but can pop-up if you're not careful.

```
2+3*sqrt(-1) #doesn't work
```

```
## Warning in sqrt(-1): NaNs produced
```

```
## [1] NaN
```

```
2+3i #works
```

```
## [1] 2+3i
```

```
class(2+3i) #complex
```

```
## [1] "complex"
```

# Assignment

In R, we can use either `←` or `=` to handle assignment.<sup>1</sup>

## Assignment with `←`

`←` is normally read aloud as "gets". You can think of it as a (left-facing) arrow saying *assign in this direction*.

```
a ← 10 + 5  
a
```

```
## [1] 15
```

Of course, an arrow can point in the other direction too (i.e. `→`). So, the following code chunk is equivalent to the previous one, although used much less frequently.

```
10 + 5 → a
```

<sup>1</sup> The `←` is really a `<` followed by a `-`. It just looks like one thing b/c of the font I'm using here.

# Assignment (cont.)

## Assignment with `=`

You can also use `=` for assignment.

```
b = 10 + 10 ## Note that the assigned object *must* be on the left with "=".  
b
```

```
## [1] 20
```

## Which assignment operator to use?

Most R users (purists?) seem to prefer `←` for assignment, since `=` also has specific role for evaluation *within* functions.

- We'll see lots of examples of this later.
- But I don't think it matters; `=` is quicker to type and is more intuitive if you're coming from another programming language. (More discussion [here](#) and [here](#).)
- I prefer `=` as it makes for easier to read and cleaner code.

**Bottom line:** Use whichever you prefer. Just be consistent.



# Data Structures in R

---

# Data Structures

In base `R`, there are 5 main data structures (not exhaustive).

1. atomic vector
2. list
3. matrix/array
4. data.frame
5. factors

# Data Structures (cont.)

## Vectors

- Vectors are a collection of multiple objects.
- There are two types of vectors:
  1. atomic vectors
  2. lists.
- Atomic vectors must be all of the same type of data.
- To create an atomic vector, put multiple objects within `c()` separated by commas.
  - **Never store an object as c!**
  - Will break your code and it might take hours to figure out why.
- A list is a collection of atomic vectors.
- Atomic vectors are one dimensional

```
numeric_grades = c(90,75,95,85,100,60,76)
letter_grades  = c("A-", "C", "A", "B", "A", "D", "C")
mixed_grades   = c("A", 95, "B", 85, "C", 75)
```

I said atomic vectors must be all of the same type. However, the last line ran without an error. What do you think happened?

# Data Structures (cont.)

## Vectors

```
class(numeric_grades)
```

```
## [1] "numeric"
```

```
class(letter_grades)
```

```
## [1] "character"
```

```
mixed_grades
```

```
## [1] "A" "95" "B" "85" "C" "75"
```

```
class(mixed_grades)
```

```
## [1] "character"
```

# Data Structures (cont.)

## Vectors: Attributes

- Objects in R can have attributes. Each type of data structure (and objects more generally) will have different attributes.
- Types of attributes include (but not limited to):
  1. names
  2. dimnames
  3. class

# Data Structures (cont.)

## Vectors: Attributes

```
names(numeric_grades)
```

```
## NULL
```

```
names(numeric_grades) = c("Student 1", "Student 2", "Student 3", "Student 4",  
                           "Student 5", "Student 6", "Student 7") #this is bad code!
```

```
names(numeric_grades)
```

```
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"
```

```
## [7] "Student 7"
```

```
names(numeric_grades) = NULL
```

```
names(numeric_grades) = paste("Student", 1:length(numeric_grades))
```

```
names(numeric_grades)
```

```
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"
```

```
## [7] "Student 7"
```

# Data Structures (cont.)

## Vectors: Indexing

- Indexing in `R` is very simple: it starts at 1 and count up by 1.
  - This is different than indexing in other programming languages which starts at 0.
  - Most "mathematical" languages will start at 1 e.g. `R`, `Matlab`, `Julia`.
- To index a vector, put `[i]` after the name of the vector where `i` is the *i*th position.

```
some_numbers = c(27,22,94)
some_numbers[1]
```

```
## [1] 27
```

```
some_numbers[3]
```

```
## [1] 94
```

```
some_numbers[1:2]
```

```
## [1] 27 22
```

# Data Structures (cont.)

## Vectors: Indexing

- Can use indexing to change values in an object

```
some_numbers[2]
```

```
## [1] 22
```

```
some_numbers[2] = 23  
some_numbers[2]
```

```
## [1] 23
```

- Can sequentially index

```
some_numbers[1:2][2]
```

```
## [1] 23
```



# Data Structures (cont.)

## Lists

- Lists are collections of atomic vectors.
- Since each atomic vector is still an atomic vector, within vector types have to be the same but across vectors can differ in type.

```
names(letter_grades) = paste("Student",1:length(letter_grades))
grade_list = list(names(numeric_grades),numeric_grades,letter_grades)
grade_list
```

```
## [[1]]
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"
## [7] "Student 7"
##
## [[2]]
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          90        75        95        85        100        60        76
##
## [[3]]
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          "A-"      "C"      "A"      "B"      "A"      "D"      "C"
```

# Data Structures (cont.)

## Lists: Indexing

- Indexing for lists is almost the same as vectors, you just have another layer to deal with.
- To return an element of the list as a list, use single brackets: `[i]`
- To return an element of the list as a vector, use double brackets: `[[i]]`

```
grade_list[1]
```

```
## [[1]]  
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"
```

```
grade_list[[1]]
```

```
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"
```

# Data Structures (cont.)

## Lists: Indexing

- Elements of a list can still have names

```
names(grade_list) = c("student names", "numeric grade", "letter grade")
grade_list
```

```
## $`student names`
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"
## [7] "Student 7"
##
## $`numeric grade`
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          90       75       95       85       100       60       76
##
## $`letter grade`
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          "A-"      "C"      "A"      "B"      "A"      "D"      "C"
```

# Data Structures (cont.)

## Lists: Indexing

- You can access an element of a list using its name and the `$`

```
grade_list$`numeric grade` #`` are used because of the space
```

```
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          90       75       95       85       100       60       76
```

- Double indexing still works

```
grade_list[[1]][4]
```

```
## [1] "Student 4"
```

```
grade_list$`student names`[4]
```

```
## [1] "Student 4"
```

# Data Structures (cont.)

## Matrix

- Matrices are essentially two dimensional atomic vectors.
- Most things that apply to atomic vectors are true for matrices
- While we can make character matrices, they aren't very useful as matrices are most useful for mathematical operations.
- Due to the nature of this course, we will probably not use matrices much as we are not coding things by hand.

```
num_mat = matrix(1:9,ncol=3)
num_mat
```

```
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
```

# Data Structures (cont.)

## Matrix: Indexing

- Indexing a matrix is similar to atomic vectors; two dimensions to index.
- If one dimension is omitted, returns that entire dimension.
- First index is for rows; second index is for columns.

```
num_mat[1,2]
```

```
## [1] 4
```

```
num_mat[1,]
```

```
## [1] 1 4 7
```

```
num_mat[,2]
```

```
## [1] 4 5 6
```

```
num_mat[5]
```

```
## [1] 5
```

# Data Structures (cont.)

## Array

- Arrays are 2 dimensional *or greater* matrices.
- Again, while arrays can be useful, we will likely not be working with them much.

```
num_array = array(1:(3*3*2),dim = c(3,3,2))
num_array
```

```
## , , 1
##
##      [,1] [,2] [,3]
## [1,]     1     4     7
## [2,]     2     5     8
## [3,]     3     6     9
##
## , , 2
##
##      [,1] [,2] [,3]
## [1,]    10    13    16
## [2,]    11    14    17
## [3,]    12    15    18
```

# Data Structures (cont.)

## Array: Indexing

- Indexing is same as matrices, just the same number of indexes as dimensions

```
num_array[1,2,1]
```

```
## [1] 4
```

```
num_array[:,1]
```

```
##      [,1] [,2] [,3]  
## [1,]    1    4    7  
## [2,]    2    5    8  
## [3,]    3    6    9
```

```
num_array[1,2,]
```

```
## [1] 4 13
```

```
num_array[12]
```

```
## [1] 12
```



# Data Structures (cont.)

## Data.frame

- data.frames are basically "data sets."
- Using established terminology, they are collections of atomic vectors in a rectangular format.
  - This may sound similar to lists; however, lists don't have a restriction on the geometric format.
  - Could have a lists of lists of lists etc.
- Each column of a data.frame must be of the same type.
- Each row is an observation.

```
data(mtcars)
head(mtcars,5)
```

```
##           mpg cyl  disp  hp  drat    wt  qsec vs  am  gear  carb
## Mazda RX4      21.0   6  160  110  3.90  2.620  16.46  0   1    4    4
## Mazda RX4 Wag  21.0   6  160  110  3.90  2.875  17.02  0   1    4    4
## Datsun 710      22.8   4  108   93  3.85  2.320  18.61  1   1    4    1
## Hornet 4 Drive  21.4   6  258  110  3.08  3.215  19.44  1   0    3    1
## Hornet Sportabout 18.7   8  360  175  3.15  3.440  17.02  0   0    3    2
```

# Data Structures (cont.)

## Data.frame

- `names()` on a `data.frame` returns the names of the variables.
- `length()` and `ncol()` return the number of variables in the `data.frame`
- `nrow()` returns the number of rows in the `data.frame`.
- Note: if you don't use `head()`, `tail()` or another way to restrict which rows are returned, returning a `data.frame` will return all rows, which is annoying.
  - This is fixed in a popular library called `data.table` that we will learn later.

`mtcars`

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

# Data Structures (cont.)

## Data.frame: Indexing

- Indexing a data.frame is similar to indexing a matrix; the columns can have names

```
mtcars[1:5, "mpg"]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7
```

```
mtcars$mpg[1:5]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7
```

```
mtcars[1:5, 1]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7
```

```
row.names(mtcars)[1:5]
```

```
## [1] "Mazda RX4"          "Mazda RX4 Wag"      "Datsun 710"  
## [4] "Hornet 4 Drive"     "Hornet Sportabout"
```

# Data Structures (cont.)

## Factors:

- Factors are essentially categorical variables.
- Factors can organized text into different groups.
- Are only really useful to data analysis and modeling.
- Factors have different "levels" which are the groups.
- Can be tricky when first learning, but ultimately rather simple.

```
fact_groups = letters[sample(1:26,10,replace=T)]  
fact_groups
```

```
## [1] "s" "o" "j" "p" "x" "j" "f" "s" "h" "s"
```

```
fact_groups = factor(fact_groups,levels=letters)  
fact_groups
```

```
## [1] s o j p x j f s h s
```

```
## Levels: a b c d e f g h i j k l m n o p q r s t u v w x y z
```

# Getting help

---

# Help

For more information on a (named) function or object in R, consult the "help" documentation. For example:

```
help(plot)
```

Or, more simply, just use `?:`

```
# This is what most people use.  
?plot
```

**Aside 1:** Comments in R are demarcated by `#`.

- Hit `Ctrl+Shift+c` in RStudio to (un)comment whole sections of highlighted code.

**Aside 2:** See the *Examples* section at the bottom of the help file?

- You can run them with the `example()` function. Try it: `example(plot)`.

# Help (cont.)

## Vignettes

For many packages, you can also try the `vignette()` function, which will provide an introduction to a package and its purpose through a series of helpful examples.

- Try running `vignette("dplyr")` in your console now.

I highly encourage reading package vignettes if they are available.

- They are often the best way to learn how to use a package.

One complication is that you need to know the exact name of the package vignette(s).

- E.g. The `dplyr` package actually has several vignettes associated with it: "dplyr", "window-functions", "programming", etc.
- You can run `vignette()` (i.e. without any arguments) to list the available vignettes of every *installed* package installed on your system.
- Or, run `vignette(all = FALSE)` if you only want to see the vignettes of any *loaded* packages.

# Help (cont.)

## Demos

Similar to vignettes, many packages come with built-in, interactive demos.

To list all available demos on your system:<sup>1</sup>

```
demo(package = .packages(all.available = TRUE))
```

To run a specific demo, just tell R which one and the name of the parent package. For example:

```
demo("graphics", package = "graphics")
```

<sup>1</sup> How would you limit the demos to one particular package?



Next lecture(s): Objects and the OOP  
approach

---