

Data Science for Economists

Lecture 2: R language basics

Jiaxi Li

University of North Carolina | ECON 370

Table of contents

1. Prologue
2. Introduction
3. Types of Data
4. Data Structures
5. Getting Help

* Slides adapted from Drew Van Kuiken and Alex Marsh's ECON 370 and originally based on materials by Grant McDermott's EC 607 at University of Oregon.

Prologue

Agenda

Today and the next lecture are going to be very hands on.

- I'll have slides as per usual, but we're going to spend a lot of time live coding together.

This is deliberate.

- I want you to get comfortable typing R commands yourself — and navigating the RStudio IDE — without resorting to copy+paste.
- Slightly more painful in the beginning, but much better payoff in the long-run.

Introduction

(Some important R concepts)

Basic Arithmetic

R is a powerful calculator and recognizes all of the standard arithmetic operators:

```
1+2 ## Addition
```

```
## [1] 3
```

```
6-7 ## Subtraction
```

```
## [1] -1
```

```
5/2 ## Division
```

```
## [1] 2.5
```

```
2^3 ## Exponentiation
```

```
## [1] 8
```

```
2+4*1^3 ## Standard order of precedence (`*` before `+`, etc.)
```

```
## [1] 6
```

Basic Arithmetic

When possible, do operators in vectors in `R`.

```
first_vec = 1:5  #store our first vector  
first_vec      #show vector
```

```
## [1] 1 2 3 4 5
```

```
first_vec + 0.5  #add 1/2 to each element
```

```
## [1] 1.5 2.5 3.5 4.5 5.5
```

```
first_vec + 6:10 #add another vector to first_vector
```

```
## [1]  7  9 11 13 15
```

```
first_vec + 6:9  #oops!
```

```
## Warning in first_vec + 6:9: longer object length is not a multiple of shorter  
## object length
```

```
## [1]  7  9 11 13 11
```

Basic Arithmetic (cont.)

We can also invoke modulo operators (quotient & remainder).

- Very useful when dealing with time, for example.

```
100 %/% 60 ## How many whole hours in 100 minutes?
```

```
## [1] 1
```

```
120 %/% 60 ## How many whole hours in 120 minutes?
```

```
## [1] 2
```

```
100 %% 60 ## How many minutes are left over from dividing 100 by 60?
```

```
## [1] 40
```

```
120 %% 60 ## How many minutes are left over from dividing 120 by 60?
```

```
## [1] 0
```


More about Vector Operation

`:` sequence: start from the first number, add 1 each time, stop right before become greater than the later number.

```
1.4:3.5
```

```
## [1] 1.4 2.4 3.4
```

You can also use `c`, concatenation to create a vector (more about `c` later):

```
c(1,2,5,12,1.4)
```

```
## [1] 1.0 2.0 5.0 12.0 1.4
```

```
c("Blue", "Beige", "Black", "Brown")
```

```
## [1] "Blue" "Beige" "Black" "Brown"
```

More about Vector Operation

In vector operation, we can have number first and then vector:

```
2^(1:4)
```

```
## [1] 2 4 8 16
```

```
1 + 2:4
```

```
## [1] 3 4 5
```

Modulo operators can also be applied to vectors:

```
110:130 %% 60
```

```
## [1] 50 51 52 53 54 55 56 57 58 59 0 1 2 3 4 5 6 7 8 9 10
```

```
120 %% 15:20
```

```
## [1] 0 8 12 6 0
```

Data Types in R

Data Types

R has 6 basic data types:

1. Character
2. Numeric
3. Integer
4. Logical
5. Complex
6. Raw (we will mostly ignore this type)

Data Types (Character)

- The character type can be described as "text."
 - It is known as a "string" in other programming languages.

```
my_name      = "Jiaxi Li"
first_name   = "Jiaxi"
last_name    = "Li"      #good style to line up to equals signs (i often don't put in the
class(my_name)
```

```
## [1] "character"
```

```
is.character(my_name)
```

```
## [1] TRUE
```

Data Types (Character)

Counting Characters

```
length(my_name) # length also works for other data types
```

```
## [1] 1
```

```
length(c(first_name,last_name))
```

```
## [1] 2
```

```
nchar(my_name)
```

```
## [1] 8
```

```
nchar(c(first_name,last_name))
```

```
## [1] 5 2
```

Data Types (Character)

Comparing (In general)

Important: `==` is used for comparing two objects and `=` is used for assigning values.

```
also_my_name = 'Jiaxi Li' # We are using ' instead of " here  
my_name = also_my_name
```

```
## [1] TRUE
```

```
Lowercase_my_name = 'jiaxi li'  
my_name = Lowercase_my_name
```

```
## [1] FALSE
```

Data Types (Character)

Combining Characters

```
paste(first_name,last_name)
```

```
## [1] "Jiaxi Li"
```

```
paste(first_name,last_name,sep="-")
```

```
## [1] "Jiaxi-Li"
```

```
paste0(first_name,last_name)
```

```
## [1] "JiaxiLi"
```

```
first_name + " " + last_name
```

```
## Error in first_name + " ": non-numeric argument to binary operator
```

Note: `c` and `paste` are different and `+` does not work for character. We can use `cat` as a combination of `paste` and `print`.

Data Types (Character)

Printing (In general)

- The following code will look *very* similar to a user when examined in the console.

```
my_name
```

```
## [1] "Jiaxi Li"
```

```
print(my_name)
```

```
## [1] "Jiaxi Li"
```

- However, what's happening is ***very different***.
 - The first is called auto-printing, and it only happens in interactive mode (e.g., console or top-level script execution).
 - The second forces printing even in non-interactive contexts (loops and functions). We will talk about loops and functions later in more details.

Data Types (Character)

Printing (In general)

Things inside of functions/curly brackets are considered non-interactive contexts

```
nchar(my_name)
```

```
## [1] 8
```

```
nchar(print(my_name))
```

```
## [1] "Jiaxi Li"
```

```
## [1] 8
```

```
for (j in 1) {my_name}  
for (j in 1) {print(my_name)}
```

```
## [1] "Jiaxi Li"
```

Data Types (Character)

Change Cases

```
toupper(paste(first_name,last_name))
```

```
## [1] "JIAXI LI"
```

```
tolower(paste(first_name,last_name))
```

```
## [1] "jiaxi li"
```

```
toupper(c(first_name,last_name))
```

```
## [1] "JIAXI" "LI"
```

```
tolower(c(first_name,last_name))
```

```
## [1] "jiaxi" "li"
```

Data Types (Character)

Special Characters

- Special characters "not native" to English can still be used but must be encoded so that `R` handles them correctly.
- Encoding guarantees that the computer will handle the special characters correctly.
- We will not spend much time on this; however, to read more about encoding text, [see this post](#).

In general, make sure to switch back to English keyboard before coding!!!!!!

Data Types (Numeric)

Numeric data are "numbers"

- In a lot of programming languages, there is this distinction between floats and integers. This is also true in R but to a much smaller degree.
- Unless explicitly stated as an integer, all "numbers" are numeric data.

```
my_age      = 29
my_height = 5 + 11/12
c(my_age,my_height)
```

```
## [1] 29.000000 5.916667
```

```
class(my_age)
```

```
## [1] "numeric"
```

```
class(my_height)
```

```
## [1] "numeric"
```

Data Types (Numeric)

Special Numeric Data

- There are a handful of special numeric values including Inf, -Inf and NaN
- Be careful when working with these

```
1/0
```

```
## [1] Inf
```

```
log(0)
```

```
## [1] -Inf
```

```
log(-1)
```

```
## Warning in log(-1): NaNs produced
```

```
## [1] NaN
```

Data Types (Numeric)

Integers

"Whole numbers"

- Integers in R are distinguished from numeric data by having an L after the number part.
- The distinction between numeric and integers is not that important here in R.
 - We will mostly ignore the distinction.

```
also_my_age = 29L  
class(my_age)
```

```
## [1] "numeric"
```

```
class(also_my_age)
```

```
## [1] "integer"
```

Data Types (Logical)

Logical data are either `TRUE` or `FALSE`.

- `T` and `F` are equivalent.
- Not `true` and `false` though

```
R_is_fun = TRUE  
R_is_hard = FALSE  
R_is_fun = T
```

```
## [1] TRUE
```

```
R_is_fun = true
```

```
## Error: object 'true' not found
```


Data Types (Logical)

Operations

The most common use of logical data is for testing conditions and control flow.

- `>`, `<`, `>=`, and `<=` test for greater/less than and/or equal to.
- `&` and `|` are the logical operators for "and" and "or" respectively.
 - Order of operations: `&` are evaluated before `|`.
- `!` negates a logical: `!TRUE` becomes `FALSE` and vice versa,
- To test if equal, use two equal signs `==`. Not equal `!=`.

Data Types (Logical)

Operations

```
2 > 1
```

```
## [1] TRUE
```

```
1 > 2 & 1 > 1/2
```

```
## [1] FALSE
```

```
1 > 2 | 1 > 1/2
```

```
## [1] TRUE
```

```
2 > 1 | 1 > 1/2
```

```
## [1] TRUE
```

Data Types (Logical)

Operations

```
1 > 1/2 & 1 > 2 | 3 > 2    #order of operations are important
```

```
## [1] TRUE
```

```
(1 > 1/2) & (1 > 2) | (3 > 2)    #above is equivalent to this
```

```
## [1] TRUE
```

```
1 > 2 & (1 > 1/2 | 3 > 2) #use parentheses to change the order of operation
```

```
## [1] FALSE
```

```
0.5 == 1/2
```

```
## [1] TRUE
```

```
3 != 2    #that is ! followed by a =
```

```
## [1] TRUE
```

Data Types (Logical)

Operations

`%in%` tests for containment

```
R_is_fun
```

```
## [1] TRUE
```

```
!R_is_fun
```

```
## [1] FALSE
```

```
2 %in% c(1,2,3,4)
```

```
## [1] TRUE
```

```
!(5 %in% c(1,2,3,4))
```

```
## [1] TRUE
```

Data Types (cont.)

Special Logicals

- `NA` (Not Available) is a special type of logical data
 - Cannot be compared the same way.
 - `is.na` can detect `NA` and `NaN`
- `NaN` (Not a Number) is a special type of numeric data
 - Calculation result is not a number
 - `is.nan` only works for `NaN`

```
NA + 1
```

```
## [1] NA
```

```
is.na(NA)
```

```
## [1] TRUE
```

```
is.nan(NaN)
```

```
## [1] TRUE
```

Data Types (Logical)

Operations

If "mathematical operations" are performed on logicals, they will be coerced into 1s and 0s.

```
as.numeric(c(TRUE, FALSE))
```

```
## [1] 1 0
```

```
TRUE + FALSE
```

```
## [1] 1
```

```
sum(c(TRUE, FALSE, TRUE))
```

```
## [1] 2
```

```
mean(c(TRUE, FALSE, TRUE))
```

```
## [1] 0.6666667
```

Data Types (Logical)

Comparing Floats

Must be careful when testing for equality of floating point numbers ("decimals")

```
0.1 + 0.2 == 0.3
```

```
## [1] FALSE
```

What happened?

- Finite precision: floating point numbers aren't exact due to how numbers are stored in a computer.

```
all.equal(0.1+0.2,0.3)
```

```
## [1] TRUE
```

Data Types (Logical)

You can read more about logical operators and types [here](#) and [here](#).

Data Types (Complex)

Complex data are data that have complex numbers.

- We probably won't use these much but can pop-up if you're not careful.

```
2+3*sqrt(-1) #doesn't work
```

```
## Warning in sqrt(-1): NaNs produced
```

```
## [1] NaN
```

```
i #doesn't work
```

```
## Error: object 'i' not found
```

```
2+3i #works
```

```
## [1] 2+3i
```

```
class(2+3i) #complex
```

```
## [1] "complex"
```

Assignment

In R, we can use either `←` or `=` to handle assignment.¹

Assignment with `←`

`←` is normally read aloud as "gets". You can think of it as a (left-facing) arrow saying *assign in this direction*.

```
a ← 10 + 5  
a
```

```
## [1] 15
```

Of course, an arrow can point in the other direction too (i.e. `→`). So, the following code chunk is equivalent to the previous one, although used much less frequently.

```
10 + 5 → a
```

¹ The `←` is really a `<` followed by a `-`. It just looks like one thing b/c of the font I'm using here.

Assignment (cont.)

Assignment with `=`

You can also use `=` for assignment.

```
b = 10 + 10 ## Note that the assigned object *must* be on the left with "=".  
b
```

```
## [1] 20
```

Which assignment operator to use?

Most R users (purists?) seem to prefer `←` for assignment, since `=` also has specific role for evaluation *within* functions.

- We'll see lots of examples of this later.
- But I don't think it matters; `=` is quicker to type and is more intuitive if you're coming from another programming language. (More discussion [here](#) and [here](#).)
- Slides for this class uses `=`, but I personally prefer `←`.

Bottom line: Use whichever you prefer. Just be consistent.

Data Structures in R

Data Structures

In base `R`, there are 5 main data structures (not exhaustive).

1. atomic vector
2. list
3. matrix/array
4. data.frame
5. factors

Data Structures (Vectors)

- Vectors are a collection of multiple objects.
- There are two types of vectors:
 1. atomic vectors
 2. lists.
- Atomic vectors must be all of the same type of data.
- To create an atomic vector, put multiple objects within `c()` separated by commas.
 - **Never store an object as c! In general, avoid naming objects after existing functions**
 - Will break your code and it might take hours to figure out why.
- A list is a collection of atomic vectors.
- Atomic vectors are one dimensional

```
numeric_grades = c(90,75,95,85,100,60,76)
letter_grades   = c("A-", "C", "A", "B", "A", "D", "C")
mixed_grades    = c("A", 95, "B", 85, "C", 75)
```

I said atomic vectors must be all of the same type. However, the last line ran without an error. What do you think happened?

Data Structures (Vectors)

```
class(numeric_grades)
```

```
## [1] "numeric"
```

```
class(letter_grades)
```

```
## [1] "character"
```

```
mixed_grades
```

```
## [1] "A" "95" "B" "85" "C" "75"
```

```
class(mixed_grades)
```

```
## [1] "character"
```

Data Structures (Vectors)

Attributes

- Objects in R can have attributes. Each type of data structure (and objects more generally) will have different attributes.
- Types of attributes include (but not limited to):
 1. names
 2. dimnames
 3. class

Data Structures (Vectors)

Attributes

```
names(numeric_grades)
```

```
## NULL
```

```
names(numeric_grades) = c("Student 1", "Student 2", "Student 3", "Student 4",  
                           "Student 5", "Student 6", "Student 7") #this is bad code!  
names(numeric_grades)
```

```
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"
```

```
names(numeric_grades) = NULL  
names(numeric_grades) = paste("Student", 1:length(numeric_grades))  
names(numeric_grades)
```

```
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"
```

Data Structures (Vectors)

Indexing

- Indexing in `R` is very simple: it starts at 1 and count up by 1.
 - This is different than indexing in other programming languages which starts at 0.
 - Most "mathematical" languages will start at 1 e.g. `R`, `Matlab`, `Julia`.
- To index a vector, put `[i]` after the name of the vector where `i` is the *i*th position.

```
some_numbers = c(27,22,94)
some_numbers[1]
```

```
## [1] 27
```

```
some_numbers[3]
```

```
## [1] 94
```

```
some_numbers[1:2]
```

```
## [1] 27 22
```

Data Structures (Vectors)

Indexing

- Can use indexing to change values in an object

```
some_numbers[2]
```

```
## [1] 22
```

```
some_numbers[2] = 23  
some_numbers[2]
```

```
## [1] 23
```

- Can sequentially index

```
some_numbers[1:2][2]
```

```
## [1] 23
```

Data Structures (Vectors)

Indexing

You can also use `c` to add object and `-` index to remove object

```
some_numbers
```

```
## [1] 27 23 94
```

```
c(some_numbers, 12)
```

```
## [1] 27 23 94 12
```

```
some_numbers[-3]
```

```
## [1] 27 23
```

```
some_numbers[-(2:3)]
```

```
## [1] 27
```

Data Structures (Lists)

- Lists are collections of atomic vectors.
- Since each atomic vector is still an atomic vector, within vector types have to be the same but across vectors can differ in type.

```
names(letter_grades) = paste("Student",1:length(letter_grades))
grade_list = list(names(numeric_grades),numeric_grades,letter_grades)
grade_list
```

```
## [[1]]
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"
## [7] "Student 7"
##
## [[2]]
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          90          75          95          85          100          60          76
##
## [[3]]
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          "A-"          "C"          "A"          "B"          "A"          "D"          "C"
```

Data Structures (Lists)

Indexing

- Indexing for lists is almost the same as vectors, you just have another layer to deal with.
- To return an element of the list as a list, use single brackets: `[i]`
- To return an element of the list as a vector, use double brackets: `[[i]]`

```
grade_list[1]
```

```
## [[1]]  
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"
```

```
grade_list[[1]]
```

```
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"
```

Data Structures (Lists)

Indexing

- Elements of a list can still have names

```
names(grade_list) = c("student names", "numeric grade", "letter grade")
grade_list
```

```
## $`student names`
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"
## [7] "Student 7"
##
## $`numeric grade`
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          90       75       95       85       100       60       76
##
## $`letter grade`
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          "A-"      "C"      "A"      "B"      "A"      "D"      "C"
```

Data Structures (Lists)

Indexing

- You can access an element of a list using its name and the `$`

```
grade_list$`numeric grade` #`` are used because of the space
```

```
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7
##          90       75       95       85       100       60       76
```

- Double indexing still works

```
grade_list[[1]][4]
```

```
## [1] "Student 4"
```

```
grade_list$`student names`[4]
```

```
## [1] "Student 4"
```


Data Structures (Lists)

Indexing

- You can add an element by assigning it to the list with `$` and remove an element by setting it to `NULL`

```
grade_list$`Newfeature` ← 1:10
```

```
grade_list
```

```
## $`student names`  
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"  
##  
## $`numeric grade`  
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7  
##          90       75       95       85       100       60       76  
##  
## $`letter grade`  
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7  
##          "A-"      "C"      "A"      "B"      "A"      "D"      "C"  
##  
## $Newfeature  
## [1]  1  2  3  4  5  6  7  8  9 10
```

Data Structures (Lists)

Indexing

- You can remove an element by setting it to `NULL`

```
grade_list$`Newfeature` ← NULL
```

```
grade_list
```

```
## $`student names`  
## [1] "Student 1" "Student 2" "Student 3" "Student 4" "Student 5" "Student 6"  
## [7] "Student 7"  
##  
## $`numeric grade`  
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7  
##          90        75        95        85        100        60        76  
##  
## $`letter grade`  
## Student 1 Student 2 Student 3 Student 4 Student 5 Student 6 Student 7  
##          "A-"        "C"        "A"        "B"        "A"        "D"        "C"
```

Data Structures (Matrix)

- Matrices are essentially two dimensional atomic vectors.
- Most things that apply to atomic vectors are true for matrices
- While we can make character matrices, they aren't very useful as matrices are most useful for mathematical operations.
- Due to the nature of this course, we will probably not use matrices much as we are not coding things by hand.

```
num_mat = matrix(1:9,ncol=3)
num_mat
```

```
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
```

Data Structures (Matrix)

Indexing

- Indexing a matrix is similar to atomic vectors; two dimensions to index.
- If one dimension is omitted, returns that entire dimension.
- First index is for rows; second index is for columns.

```
num_mat[1,2]
```

```
## [1] 4
```

```
num_mat[1,]
```

```
## [1] 1 4 7
```

```
num_mat[,2]
```

```
## [1] 4 5 6
```

```
num_mat[5]
```

```
## [1] 5
```

Data Structures (Matrix)

Indexing

- You can add a row using `rbind`, add a column using `cbind`.

```
cbind(num_mat, 1:3)
```

```
##      [,1] [,2] [,3] [,4]  
## [1,]    1    4    7    1  
## [2,]    2    5    8    2  
## [3,]    3    6    9    3
```

```
rbind(num_mat, 1:3)
```

```
##      [,1] [,2] [,3]  
## [1,]    1    4    7  
## [2,]    2    5    8  
## [3,]    3    6    9  
## [4,]    1    2    3
```

- You can still remove row/column using `-` (You cannot remove an single entry though).

```
num_mat[,-2]
```

Data Structures (Array)

- Arrays are 2 dimensional *or greater* matrices.
- Again, while arrays can be useful, we will likely not be working with them much.

```
num_array = array(1:(3*3*2),dim = c(3,3,2))
num_array
```

```
## , , 1
##
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
##
## , , 2
##
##      [,1] [,2] [,3]
## [1,]   10   13   16
## [2,]   11   14   17
## [3,]   12   15   18
```

Data Structures (Array)

Indexing

- Indexing is same as matrices, just the same number of indexes as dimensions

```
num_array[1,2,1]
```

```
## [1] 4
```

```
num_array[:,1]
```

```
##      [,1] [,2] [,3]  
## [1,]    1    4    7  
## [2,]    2    5    8  
## [3,]    3    6    9
```

```
num_array[1,2,]
```

```
## [1] 4 13
```

```
num_array[12]
```

```
## [1] 12
```

Data Structures (Data.frame)

- data.frames are basically "data sets."
- Using established terminology, they are collections of atomic vectors in a rectangular format.
 - This may sound similar to lists; however, lists don't have a restriction on the geometric format.
 - Could have a lists of lists of lists etc.
- Each column of a data.frame must be of the same type.
- Each row is an observation.

```
data(mtcars)
head(mtcars,5)
```

```
##           mpg cyl  disp  hp  drat    wt  qsec vs  am  gear  carb
## Mazda RX4      21.0   6  160  110  3.90  2.620  16.46  0   1    4    4
## Mazda RX4 Wag  21.0   6  160  110  3.90  2.875  17.02  0   1    4    4
## Datsun 710     22.8   4  108   93  3.85  2.320  18.61  1   1    4    1
## Hornet 4 Drive  21.4   6  258  110  3.08  3.215  19.44  1   0    3    1
## Hornet Sportabout 18.7   8  360  175  3.15  3.440  17.02  0   0    3    2
```


Data Structures (Data.frame)

- `names()` on a data.frame returns the names of the variables.
- `length()` and `ncol()` return the number of variables in the data.frame
- `nrow()` returns the number of rows in the data.frame.
- Note: if you don't use `head()`, `tail()` or another way to restrict which rows are returned, returning a data.frame will return all rows, which is annoying.
 - This is fixed in a popular library called `data.table` that we will learn later.
 - This is also fixed in `tibble`, a modern reimagining of the data frame, in `tidyverse`.

mtcars

```
##           mpg  cyl  disp  hp drat    wt  qsec vs am gear carb
## Mazda RX4      21.0    6 160.0 110 3.90 2.620 16.46 0  1    4    4
## Mazda RX4 Wag  21.0    6 160.0 110 3.90 2.875 17.02 0  1    4    4
## Datsun 710      22.8    4 108.0  93 3.85 2.320 18.61 1  1    4    1
## Hornet 4 Drive  21.4    6 258.0 110 3.08 3.215 19.44 1  0    3    1
## Hornet Sportabout 18.7    8 360.0 175 3.15 3.440 17.02 0  0    3    2
## Valiant         18.1    6 225.0 105 2.76 3.460 20.22 1  0    3    1
## Duster 360      14.3    8 360.0 245 3.21 3.570 15.84 0  0    3    4
## Merc 240D       24.4    4 146.7  62 3.69 3.190 20.00 1  0    4    2
## Merc 230        22.8    4 140.8  95 3.92 3.150 22.90 1  0    4    2
## Merc 280        19.2    6 167.6 123 3.92 3.440 18.30 1  0    4    4
## Merc 280C       17.8    6 167.6 123 3.92 3.440 18.90 1  0    4    4
```

Data Structures (Data.frame)

Indexing

- Indexing a data.frame is similar to indexing a matrix; the columns can have names

```
mtcars[1:5, "mpg"]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7
```

```
mtcars$mpg[1:5]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7
```

```
mtcars[1:5, 1]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7
```

```
row.names(mtcars)[1:5]
```

```
## [1] "Mazda RX4"           "Mazda RX4 Wag"       "Datsun 710"  
## [4] "Hornet 4 Drive"      "Hornet Sportabout"
```

Data Structures (Data.frame)

Indexing

- Adding/removing row/column is also similar to matrix using `rbind` / `cbind` / `-`. Also, we can use `$` or assing to `NULL` similar to list.

```
mtcars[-1,]  
mtcars[, -1]  
mtcars$NewVar ← 1  
mtcars$NewVar ← NULL
```

Need same number of columns when doing rbind

```
rbind(mtcars, 1:11)
```

Need same number of rows when doing cbind

```
cbind(mtcars, 1:32)
```

Data Structures (Factors)

- Factors are essentially categorical variables.
- Factors can organized text into different groups.
- Are only really useful to data analysis and modeling.
- Factors have different "levels" which are the groups.
- Can be tricky when first learning, but ultimately rather simple.

```
# Draw 10 random letters with replacement  
fact_groups = letters[sample(1:26,10,replace=T)]  
fact_groups
```

```
## [1] "m" "h" "a" "l" "z" "h" "a" "k" "v" "v"
```

```
# Form the factor group  
fact_groups = factor(fact_groups,levels=letters)  
fact_groups
```

```
## [1] m h a l z h a k v v
```

```
## Levels: a b c d e f g h i j k l m n o p q r s t u v w x y z
```

Getting help

Help

For more information on a (named) function or object in R, consult the "help" documentation. For example:

```
help(plot)
```

Or, more simply, just use `?:`

```
# This is what most people use.  
?plot
```

Aside: See the *Examples* section at the bottom of the help file?

- You can run them with the `example()` function. Try it: `example(plot)`.

Help (cont.)

Vignettes

For many packages, you can also try the `vignette()` function, which will provide an introduction to a package and its purpose through a series of helpful examples.

- Try running `vignette("dplyr")` in your console now.

I highly encourage reading package vignettes if they are available.

- They are often the best way to learn how to use a package.

One complication is that you need to know the exact name of the package vignette(s).

- E.g. The `dplyr` package actually has several vignettes associated with it: "dplyr", "window-functions", "programming", etc.
- You can run `vignette()` (i.e. without any arguments) to list the available vignettes of every *installed* package installed on your system.
- Or, run `vignette(all = FALSE)` if you only want to see the vignettes of any *loaded* packages.

Help (cont.)

Demos

Similar to vignettes, many packages come with built-in, interactive demos.

To list all available demos on your system:¹

```
demo(package = .packages(all.available = TRUE))
```

To run a specific demo, just tell R which one and the name of the parent package. For example:

```
demo("graphics", package = "graphics")
```

¹ How would you limit the demos to one particular package?

HW 1

I will release HW 1 on this afternoon

It will be due on Tuesday, 9/11, before our class begins.

...*because* this is introductory material, I am going to ask that you **do not** use LLMs/AI models for help with this homework specifically.

I know I said you could (and should!) use them for our class. And you should! But you need to build strong coding foundations first

Recommended order of resources:

Try Yourself → Slides → Online Threads → Classmates → ChatGPT (but ask for explanation)

Next lecture(s): Objects and the OOP approach
