

Data Science for Economists

Lecture 7: Data cleaning & wrangling: Tidyverse

Jiaxi Li

University of North Carolina | ECON 370

Table of contents

1. Prologue
2. Tidyverse basics
3. Data wrangling with dplyr
 - filter
 - arrange
 - select
 - mutate
 - summarize
 - joins
4. Data tidying with tidyr
 - pivot_longer / pivot_wider
 - separate and unite
5. Summary

Prologue

What is "tidy" data?

Resources:

- [Vignette](#) (from the **tidyr** package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Basically, tidy data is more likely to be [long \(i.e. narrow\) format](#) than wide format.

Checklist

R packages you'll need today

☑ tidyverse

☑ nycflights13

I'll hold off loading these libraries for now. But you can install/update them both with the following command.

```
install.packages(c('tidyverse', 'nycflights13'), repos = 'https://cran.rstudio.com', c
```

Tip: If you're on Linux, then I *strongly* recommend installing the pre-compiled binary versions of these packages from **RSPM** instead of CRAN. The exact repo mirror varies by distro (see the link). But on Ubuntu 20.04, for example, you'd use:

```
install.packages(c('tidyverse', 'nycflights13'), repos = 'https://packagemanager.rstuc
```

Tidyverse basics

Tidyverse vs. base R

Much digital ink has been spilled over the "tidyverse vs. base R" debate.

I won't delve into this debate here, because I think the answer is **clear**: We should teach the **tidyverse** first (or, at least, early).

- The documentation and community support are outstanding.
- Having a consistent philosophy and syntax makes it easier to learn.
- Provides a convenient "front-end" to big data tools that we'll use later in the course.
- For **data cleaning, wrangling, and plotting**, the tidyverse really is a no-brainer.¹

But... this certainly shouldn't put you off learning **base R** alternatives.

- Base R is extremely flexible and powerful (and stable).
- There are some things that you'll have to venture outside of the tidyverse for (**loops, functions, etc.**).
- A combination of tidyverse and base R is often the best solution to a problem.
- Excellent base R data manipulation tutorials: [here](#) and [here](#).

Tidyverse vs. base R (cont.)

One point of convenience is that there is often a direct correspondence between a tidyverse command and its base R equivalent.

These generally follow a `tidyverse::snake_case` vs `base::period.case` rule. E.g. Compare:

tidyverse	base
<code>?readr::read_csv</code>	<code>?utils::read.csv</code>
<code>?dplyr::if_else</code>	<code>?base::ifelse</code>
<code>?tibble::tibble</code>	<code>?base::data.frame</code>

Etcetera.

If you call up the above examples, you'll see that the tidyverse alternative typically offers some enhancements or other useful options (and sometimes restrictions) over its base counterpart.

Remember: There are (almost) always multiple ways to achieve a single goal in R.

Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats    1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2    3.5.2      ✓ tibble     3.3.0
## ✓ lubridate  1.9.4      ✓ tidyr      1.3.1
## ✓ purrr      1.0.4
## — Conflicts ————— tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to
```

We see that we have actually loaded a number of packages (these are **core** packages, which could also be loaded individually): **ggplot2**, **tibble**, **dplyr**, etc.

- We can also see information about the package versions and some namespace conflicts.

Tidyverse packages (cont.)

The tidyverse actually comes with a lot more packages than those that are just loaded automatically.¹

```
tidyverse_packages()
```

```
## [1] "broom"          "conflicted"    "cli"           "dbplyr"
## [5] "dplyr"          "dtplyr"        "forcats"       "ggplot2"
## [9] "googledrive"    "googlesheets4" "haven"         "hms"
## [13] "httr"           "jsonlite"      "lubridate"     "magrittr"
## [17] "modelr"         "pillar"        "purrr"         "ragg"
## [21] "readr"          "readxl"        "reprex"        "rlang"
## [25] "rstudioapi"     "rvest"         "stringr"       "tibble"
## [29] "tidyr"          "xml2"          "tidyverse"
```

- E.g. The **lubridate** package for working with dates and the **rvest** package for webscraping.
- However, bear in mind that these **optional** packages will have to be loaded separately.

¹ It also includes a *lot* of dependencies upon installation. This is a matter of some **controversy**.

Tidyverse packages (cont.)

I hope to cover most of the tidyverse packages over the length of this course.

Today, however, I'm only really going to focus on two packages:

1. **dplyr**
2. **tidyr**

These are the workhorse packages for cleaning and wrangling data. They are thus the ones that you will likely make the most use of (alongside **ggplot2**, which we already met back in Lecture 1).

- Data cleaning and wrangling occupies an inordinate amount of time, no matter where you are in your research career.

dplyr

Aside: dplyr 1.0.0 release

Some of the **dplyr** features that we'll cover today were introduced in **version 1.0.0** of the package.

- Version 1.0.0 is a big deal since it marks a stable code base for the package going forward.
- Please make sure that you are running at least **dplyr** 1.0.0 before continuing.

```
packageVersion('dplyr')
```

```
## [1] '1.1.4'
```

```
# install.packages('dplyr') ## install updated version if < 1.0.0
```

Note: **dplyr** 1.0.0 also notifies you about grouping variables every time you do operations on or with them. YMMV, but, personally, I find these messages annoying. You can **switch them off**.

```
options(dplyr.summarize.inform = FALSE) ## Add to .Rprofile to make permanent
```

Key dplyr verbs

There are five key dplyr verbs that you need to learn.

1. `filter`: Filter (i.e. subset) rows based on their values.
2. `arrange`: Arrange (i.e. reorder) rows based on their values.
3. `select`: Select (i.e. subset) columns by their names:
4. `mutate`: Create new columns.
5. `summarize`: Collapse multiple rows into a single summary value.¹

Let's practice these commands together using the `starwars` data frame that comes pre-packaged with dplyr.

¹ `summarise` with a "s" works too.

1) dplyr::filter

We can chain multiple filter commands with the pipe (`>`), or just separate them within a single filter command using commas.

```
starwars >
  filter(
    species = "Human",
    height ≥ 190
  )
```

```
## # A tibble: 4 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Darth Va...   202   136 none       white      yellow      41.9 male  masculi
## 2 Qui-Gon ...   193    89 brown      fair       blue        92   male  masculi
## 3 Dooku        193    80 white      fair       brown      102   male  masculi
## 4 Bail Pre...   191   NA black      tan        brown       67   male  masculi
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

What's the base R equivalent of this code?

1) dplyr::filter cont.

Regular expressions work well too.

```
starwars >
  filter(grepl("Skywalker", name))
```

```
## # A tibble: 3 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Luke Sky...   172    77 blond      fair        blue         19  male  mascu...
## 2 Anakin S...   188    84 blond      fair        blue        41.9  male  mascu...
## 3 Shmi Sky...   163    NA black      fair        brown        72  fema... femin...
```

i 5 more variables: homeworld <chr>, species <chr>, films <list>,
vehicles <list>, starships <list>

1) dplyr::filter cont.

A very common `filter` use case is identifying (or removing) missing data cases.

```
starwars >
  filter(is.na(height))
```

```
## # A tibble: 6 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Arvel Cr...    NA    NA brown      fair        brown          NA male  mascu...
## 2 Finn           NA    NA black      dark        dark           NA male  mascu...
## 3 Rey            NA    NA brown      light       hazel           NA fema... femin...
## 4 Poe Dame...    NA    NA brown      light       brown           NA male  mascu...
## 5 BB8            NA    NA none       none        black           NA none  mascu...
## 6 Captain ...    NA    NA none       none        unknown         NA fema... femin...
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

To remove missing observations, simply use negation: `filter(!is.na(height))`. Try this yourself.

2) dplyr::arrange

```
starwars ▶  
  arrange(birth_year)
```

```
## # A tibble: 87 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Wicket ...    88   20  brown      brown      brown         8  male  masculi  
## 2 IG-88        200  140  none       metal      red          15  none  masculi  
## 3 Luke Sk...   172   77  blond      fair       blue         19  male  masculi  
## 4 Leia Or...   150   49  brown      light      brown        19  fema... femin...  
## 5 Wedge A...   170   77  brown      fair       hazel        21  male  masculi  
## 6 Plo Koon     188   80  none       orange     black        22  male  masculi  
## 7 Biggs D...   183   84  black      light      brown        24  male  masculi  
## 8 Han Solo     180   80  brown      fair       brown        29  male  masculi  
## 9 Lando C...   177   79  black      dark       brown        31  male  masculi  
## 10 Boba Fe...   183  78.2  black      fair       brown       31.5  male  masculi  
## # i 77 more rows  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

Note: Arranging on a character-based column (i.e. strings) will sort alphabetically. Try this yourself by arranging according to the "name" column.

2) dplyr::arrange cont.

We can also arrange items in descending order using `arrange(desc())`.

```
starwars ►  
  arrange(desc(birth_year))
```

```
## # A tibble: 87 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Yoda          66    17 white      green      brown          896 male  masculi  
## 2 Jabba D...    175   1358 <NA>      green-tan... orange          600 herm... masculi  
## 3 Chewbac...    228    112 brown     unknown    blue           200 male  masculi  
## 4 C-3PO        167     75 <NA>      gold       yellow          112 none  masculi  
## 5 Dooku         193     80 white     fair       brown          102 male  masculi  
## 6 Qui-Gon...    193     89 brown     fair       blue            92 male  masculi  
## 7 Ki-Adi-...    198     82 white     pale      yellow          92 male  masculi  
## 8 Finis V...    170     NA blond     fair       blue            91 male  masculi  
## 9 Palpati...    170     75 grey     pale      yellow          82 male  masculi  
## 10 Cliegg ...    183     NA brown     fair       blue            82 male  masculi  
## # i 77 more rows  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

3) dplyr::select

Use commas to select multiple columns out of a data frame. (You can also use "first:last" for consecutive columns). Deselect a column with "-".

```
starwars ►  
  select(name:skin_color, species, -height)
```

```
## # A tibble: 87 × 5  
##   name                mass hair_color  skin_color species  
##   <chr>              <dbl> <chr>      <chr>      <chr>  
## 1 Luke Skywalker      77 blond      fair        Human  
## 2 C-3PO                75 <NA>      gold        Droid  
## 3 R2-D2                32 <NA>      white, blue Droid  
## 4 Darth Vader         136 none      white        Human  
## 5 Leia Organa          49 brown      light        Human  
## 6 Owen Lars           120 brown, grey light        Human  
## 7 Beru Whitesun Lars   75 brown      light        Human  
## 8 R5-D4                32 <NA>      white, red  Droid  
## 9 Biggs Darklighter   84 black      light        Human  
## 10 Obi-Wan Kenobi      77 auburn, white fair         Human  
## # i 77 more rows
```

3) dplyr::select *cont.*

You can also rename some (or all) of your selected variables in place.

```
starwars ►  
  select(alias=name, crib=homeworld, sex=gender)
```

```
## # A tibble: 87 × 3  
##   alias          crib      sex  
##   <chr>         <chr>   <chr>  
## 1 Luke Skywalker Tatooine masculine  
## 2 C-3PO          Tatooine masculine  
## 3 R2-D2          Naboo    masculine  
## 4 Darth Vader    Tatooine masculine  
## 5 Leia Organa    Alderaan feminine  
## 6 Owen Lars      Tatooine masculine  
## 7 Beru Whitesun Lars Tatooine feminine  
## 8 R5-D4          Tatooine masculine  
## 9 Biggs Darklighter Tatooine masculine  
## 10 Obi-Wan Kenobi Stewjon  masculine  
## # i 77 more rows
```

If you just want to rename columns without subsetting them, you can use `rename`. Try this now by replacing `select( ... )` in the above code chunk with `rename( ... )`.

3) dplyr::select cont.

The `select(contains(PATTERN))` option provides a nice shortcut in relevant cases.

```
starwars >
  select(name, contains("color"))
```



```
## # A tibble: 87 × 4
```

	name	hair_color	skin_color	eye_color
	<chr>	<chr>	<chr>	<chr>
## 1	Luke Skywalker	blond	fair	blue
## 2	C-3PO	<NA>	gold	yellow
## 3	R2-D2	<NA>	white, blue	red
## 4	Darth Vader	none	white	yellow
## 5	Leia Organa	brown	light	brown
## 6	Owen Lars	brown, grey	light	blue
## 7	Beru Whitesun Lars	brown	light	blue
## 8	R5-D4	<NA>	white, red	red
## 9	Biggs Darklighter	black	light	brown
## 10	Obi-Wan Kenobi	auburn, white	fair	blue-gray

```
## # i 77 more rows
```

There are more functions can be used with `select` such as `start_with()`, `end_with()`, `matches()` and `num_range()`.

3) dplyr::select cont.

The `select(..., everything())` option is another useful shortcut if you only want to bring some variable(s) to the "front" of a data frame.

```
starwars >
  select(species, homeworld, everything()) >
  head(5)
```

A tibble: 5 × 14

	species	homeworld	name	height	mass	hair_color	skin_color	eye_color
	<chr>	<chr>	<chr>	<int>	<dbl>	<chr>	<chr>	<chr>
## 1	Human	Tatooine	Luke Skywalker	172	77	blond	fair	blue
## 2	Droid	Tatooine	C-3PO	167	75	<NA>	gold	yellow
## 3	Droid	Naboo	R2-D2	96	32	<NA>	white, blue	red
## 4	Human	Tatooine	Darth Vader	202	136	none	white	yellow
## 5	Human	Alderaan	Leia Organa	150	49	brown	light	brown

i 6 more variables: birth_year <dbl>, sex <chr>, gender <chr>, films <list>,
vehicles <list>, starships <list>

Note: The `relocate` function available in dplyr 1.0.0 has brought a lot more functionality to ordering of columns. See [here](#).

4) dplyr::mutate

You can create new columns from scratch, or (more commonly) as transformations of existing columns.

```
starwars >
  select(name, birth_year) >
  mutate(dog_years = birth_year * 7) >
  mutate(comment = paste0(name, " is ", dog_years, " in dog years."))
```

```
## # A tibble: 87 × 4
##   name                birth_year dog_years comment
##   <chr>                <dbl>     <dbl> <chr>
## 1 Luke Skywalker         19        133 Luke Skywalker is 133 in dog years.
## 2 C-3PO                 112        784 C-3PO is 784 in dog years.
## 3 R2-D2                  33        231 R2-D2 is 231 in dog years.
## 4 Darth Vader           41.9       293.3 Darth Vader is 293.3 in dog years.
## 5 Leia Organa           19        133 Leia Organa is 133 in dog years.
## 6 Owen Lars             52        364 Owen Lars is 364 in dog years.
## 7 Beru Whitesun Lars    47        329 Beru Whitesun Lars is 329 in dog yea...
## 8 R5-D4                 NA         NA R5-D4 is NA in dog years.
## 9 Biggs Darklighter     24        168 Biggs Darklighter is 168 in dog year...
## 10 Obi-Wan Kenobi       57        399 Obi-Wan Kenobi is 399 in dog years.
## # i 77 more rows
```


4) dplyr::mutate cont.

Note: `mutate` is order aware. So you can chain multiple mutates in a single call.

```
starwars >
  select(name, birth_year) >
  mutate(
    dog_years = birth_year * 7, ## Separate with a comma
    comment = paste0(name, " is ", dog_years, " in dog years.")
  )
```

```
## # A tibble: 87 × 4
```

```
##   name                birth_year dog_years comment
##   <chr>                <dbl>     <dbl> <chr>
## 1 Luke Skywalker        19         133 Luke Skywalker is 133 in dog years.
## 2 C-3PO                 112         784 C-3PO is 784 in dog years.
## 3 R2-D2                  33         231 R2-D2 is 231 in dog years.
## 4 Darth Vader           41.9        293.3 Darth Vader is 293.3 in dog years.
## 5 Leia Organa           19         133 Leia Organa is 133 in dog years.
## 6 Owen Lars              52         364 Owen Lars is 364 in dog years.
## 7 Beru Whitesun Lars     47         329 Beru Whitesun Lars is 329 in dog yea...
## 8 R5-D4                  NA          NA  R5-D4 is NA in dog years.
## 9 Biggs Darklighter     24         168 Biggs Darklighter is 168 in dog year...
## 10 Obi-Wan Kenobi        57         399 Obi-Wan Kenobi is 399 in dog years.
## # i 77 more rows
```

4) dplyr::mutate cont.

Boolean, logical and conditional operators all work well with `mutate` too.

```
starwars >
  select(name, height) >
  filter(name %in% c("Luke Skywalker", "Anakin Skywalker")) >
  mutate(tall1 = height > 180) >
  mutate(tall2 = ifelse(height > 180, "Tall", "Short")) ## Same effect, but can choose
```

```
## # A tibble: 2 × 4
##   name          height tall1 tall2
##   <chr>         <int> <lgl> <chr>
## 1 Luke Skywalker    172 FALSE Short
## 2 Anakin Skywalker    188  TRUE  Tall
```

4) dplyr::mutate *cont.*

Lastly, combining `mutate` with the `across` feature in dplyr 1.0.0+ allows you to easily work on a subset of variables. For example:

```
starwars >
  select(name:eye_color) >
  mutate(across(where(is.character), toupper)) >
  head(5)
```

```
## # A tibble: 5 × 6
##   name          height  mass hair_color skin_color eye_color
##   <chr>         <int> <dbl> <chr>      <chr>    <chr>
## 1 LUKE SKYWALKER   172    77 BLOND      FAIR     BLUE
## 2 C-3PO           167    75 <NA>      GOLD     YELLOW
## 3 R2-D2           96     32 <NA>      WHITE,   RED
## 4 DARTH VADER     202   136 NONE      WHITE    YELLOW
## 5 LEIA ORGANA     150    49 BROWN     LIGHT    BROWN
```

Note: This workflow (i.e. combining `mutate` and `across`) supersedes the old "scoped" variants of `mutate` that you might have used previously. More details [here](#) and [here](#).

5) dplyr::summarize

Particularly useful in combination with the `group_by` command.

```
starwars >
  group_by(species, gender) >
  summarize(mean_height = mean(height, na.rm = TRUE), .groups = "drop")
```

```
## # A tibble: 42 × 3
##   species    gender mean_height
##   <chr>      <chr>      <dbl>
## 1 Aleena    masculine      79
## 2 Besalisk  masculine     198
## 3 Cerean    masculine     198
## 4 Chagrian  masculine     196
## 5 Clawdite  feminine     168
## 6 Droid     feminine      96
## 7 Droid     masculine    140
## 8 Dug       masculine    112
## 9 Ewok      masculine     88
## 10 Geonosian masculine    183
## # i 32 more rows
```

What would happen if we remove "na.rm = TRUE"? What does ".groups = 'drop'" do?

5) dplyr::summarize *cont.*

The same `across`-based workflow that we saw with `mutate` a few slides back also works with `summarize`. For example:

```
starwars >
  group_by(species) >
  summarize(across(where(is.numeric), function(x) mean(x, na.rm=T))) >
  head(5)
```

```
## # A tibble: 5 × 4
##   species height  mass birth_year
##   <chr>    <dbl> <dbl>      <dbl>
## 1 Aleena      79     15         NaN
## 2 Besalisk   198    102         NaN
## 3 Cerean     198     82          92
## 4 Chagrian   196    NaN         NaN
## 5 Clawdite   168     55         NaN
```

Note: Again, this functionality supersedes the old "scoped" variants of `summarize` that you used prior to dplyr 1.0.0. Details [here](#) and [here](#).

Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping.

- Particularly useful with the `summarize` and `mutate` commands, as we've already seen.

`slice`: Subset rows by position rather than filtering by values.

- E.g. `starwars > slice(c(1, 5))`

`pull`: Extract a column from a data frame as a vector or scalar.

- E.g. `starwars > filter(gender="female") > pull(height)`

`count`, `n_distinct` and `distinct`: Number and isolate unique observations.

- E.g. `starwars > count(species)`, or `starwars > n_distinct(species)`, or `starwars > distinct(species)`
- You could also use a combination of `mutate`, `group_by`, and `n()`, e.g. `starwars > group_by(species) > mutate(num = n())`.

Other dplyr goodies (cont.)

There are also a whole class of **window functions** for getting leads and lags, ranking, creating cumulative aggregates, etc.

- See `vignette("window-functions")`.

The final set of dplyr "goodies" are the family of join operations. However, these are important enough that I want to go over some concepts in a bit more depth...


- We will encounter and practice these many more times as the course progresses.

Joins

One of the mainstays of the dplyr package is merging data with the family **join operations**.

- `inner_join(df1, df2)`
- `left_join(df1, df2)`
- `right_join(df1, df2)`
- `full_join(df1, df2)`
- `semi_join(df1, df2)`
- `anti_join(df1, df2)`

Visualizing Joins



For the simple examples that I'm going to show here, we'll need some data sets that come bundled with the **nycflights13** package.

- Load it now and then inspect these data frames in your own console.

```
library(nycflights13)
flights
planes
```

Joins (cont.)

Let's perform a **left join** on the flights and planes datasets.

- *Note:* I'm going subset columns after the join, but only to keep text on the slide.

```
left_join(flights, planes) ▷  
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
```

```
## Joining with by = join_by(year, tailnum)
```

```
## # A tibble: 336,776 × 10
```

```
##   year month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>   <chr> <chr>  
## 1  2013     1     1     517     830 UA      1545 N14228 <NA> <NA>  
## 2  2013     1     1     533     850 UA      1714 N24211 <NA> <NA>  
## 3  2013     1     1     542     923 AA      1141 N619AA <NA> <NA>  
## 4  2013     1     1     544    1004 B6       725 N804JB <NA> <NA>  
## 5  2013     1     1     554     812 DL       461 N668DN <NA> <NA>  
## 6  2013     1     1     554     740 UA      1696 N39463 <NA> <NA>  
## 7  2013     1     1     555     913 B6       507 N516JB <NA> <NA>  
## 8  2013     1     1     557     709 EV      5708 N829AS <NA> <NA>  
## 9  2013     1     1     557     838 B6        79 N593JB <NA> <NA>  
## 10 2013     1     1     558     753 AA       301 N3ALAA <NA> <NA>
```

```
## # i 336,766 more rows
```

Joins (cont.)

(continued from previous slide)

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name). It also told us its choices:

```
## Joining, by = c("year", "tailnum")
```

However, there's an obvious problem here: the variable "year" does not have a consistent meaning across our joining datasets!

- In one it refers to the *year of flight*, in the other it refers to *year of construction*.

Luckily, there's an easy way to avoid this problem.

- See if you can figure it out before turning to the next slide.
- Try `?dplyr::join`.

Joins (cont.)

(continued from previous slide)

You just need to be more explicit in your join call by using the `by =` argument.

- You can also rename any ambiguous columns to avoid confusion.

```
left_join(
  flights > rename(flight_year = year),
  planes > rename(year_built = year), ## Not necessary w/ below line, but helpful
  by = "tailnum" ## Be specific about the joining column
) >
select(flight_year, month, day, dep_time, arr_time, carrier, flight, tailnum, year_built)
head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 × 11
##   flight_year month   day dep_time arr_time carrier flight tailnum year_built
##         <int> <int> <int>   <int>   <int> <chr>   <int> <chr>      <int>
## 1         2013     1     1     517     830 UA       1545 N14228     1999
## 2         2013     1     1     533     850 UA       1714 N24211     1998
## 3         2013     1     1     542     923 AA       1141 N619AA     1990
## # i 2 more variables: type <chr>, model <chr>
```

Joins (cont.)

(continued from previous slide)

Last thing I'll mention for now; note what happens if we again specify the join column... but don't rename the ambiguous "year" column in at least one of the given data frames.

```
left_join(
  flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum"
) ▷
select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, 1)
head(3)
```

```
## # A tibble: 3 × 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type  model
##   <int>  <int> <int> <int>   <int>   <int> <chr>    <int> <chr>  <chr> <chr>
## 1  2013   1999     1     1     517     830 UA       1545 N14228  Fixe... 737-...
## 2  2013   1998     1     1     533     850 UA       1714 N24211  Fixe... 737-...
## 3  2013   1990     1     1     542     923 AA       1141 N619AA  Fixe... 757-...
```

Make sure you know what "year.x" and "year.y" are. Again, it pays to be specific.

Joins (cont.)

When joining by one variable but with different names in different data

```
Tampa_flights = left_join(  
  flights,  
  airports,  
  by=join_by(dest = faa)  
) ▷  
filter(name = "Tampa Intl")
```

The alternative is to rename before joining

```
Tampa_flights = left_join(  
  flights,  
  airports ▷ rename(dest = faa),  
  by=join_by(dest)  
) ▷  
filter(name = "Tampa Intl")
```

Practice with joins

Now, please take one minute to work on the following question.

For Tampa Intl airport (Tampa Intl) on Jan 1, 2013, figure out:

Q1. Which carrier flew the most flights into Tampa?

Q2. How many Airbus (AIRBUS) planes did JetBlue (JetBlue Airways) fly into Tampa?

Q3. What manufacturers did Delta (Delta Air Lines Inc.) fly into Tampa?

tidyr

Key tidyr verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt").¹
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast").²
3. `separate`: Separate (i.e. split) one column into multiple columns.
4. `unite`: Unite (i.e. combine) multiple columns into one.

Let's practice these verbs together in class.

- Side question: Which of `pivot_longer` vs `pivot_wider` produces "tidy" data?

¹ Updated version of `tidyr::gather`.

² Updated version of `tidyr::spread`.

1) tidyr::pivot_longer

```
stocks = data.frame( ## Could use "tibble" instead of "data.frame" if you prefer
  time = as.Date('2009-01-01') + 0:1,
  X = rnorm(2, 0, 1),
  Y = rnorm(2, 0, 2),
  Z = rnorm(2, 0, 4))
stocks
```

```
##           time           X           Y           Z
## 1 2009-01-01 0.01341152 1.644425 7.199357
## 2 2009-01-02 1.23304916 1.506199 3.017213
```

```
tidy_stocks = stocks > pivot_longer(-time, names_to="stock", values_to="price")
tidy_stocks
```

```
## # A tibble: 6 × 3
##   time      stock price
##   <date>   <chr> <dbl>
## 1 2009-01-01 X     0.0134
## 2 2009-01-01 Y     1.64
## 3 2009-01-01 Z     7.20
## 4 2009-01-02 X     1.23
## 5 2009-01-02 Y     1.51
## 6 2009-01-02 Z     3.02
```

1) tidyr::pivot_longer cont.

We can also pivot_longer and separate multiple variables. More info can be found [here](#).

```
stocks2
```

```
##           time      X_price  Y_price  Z_price X_vol Y_vol Z_vol
## 1 2009-01-01 0.01341152 1.644425 7.199357      1      2      4
## 2 2009-01-02 1.23304916 1.506199 3.017213      1      2      4
```

```
stocks2_long = stocks2 >
  pivot_longer(-time,
               names_sep = "_",
               names_to = c("stock", ".value"))
```

```
stocks2_long
```

```
## # A tibble: 6 × 4
##   time      stock price  vol
##   <date>    <chr> <dbl> <dbl>
## 1 2009-01-01 X      0.0134      1
## 2 2009-01-01 Y      1.64        2
## 3 2009-01-01 Z      7.20        4
## 4 2009-01-02 X      1.23        1
## 5 2009-01-02 Y      1.51        2
## 6 2009-01-02 Z      3.02        4
```

2) tidyr::pivot_wider

```
tidy_stocks > pivot_wider(names_from=stock, values_from=price)
```

```
## # A tibble: 2 × 4  
##   time          X      Y      Z  
##   <date>      <dbl> <dbl> <dbl>  
## 1 2009-01-01 -1.13  -1.55  -4.51  
## 2 2009-01-02  0.381 -0.923  2.99
```

```
tidy_stocks > pivot_wider(names_from=time, values_from=price)
```

```
## # A tibble: 3 × 3  
##   stock 2009-01-01 2009-01-02  
##   <chr>      <dbl>      <dbl>  
## 1 X          -1.13         0.381  
## 2 Y          -1.55        -0.923  
## 3 Z          -4.51         2.99
```

Note that the second example — which has combined different pivoting arguments — has effectively transposed the data.

Aside: Remembering the `pivot_*` syntax

There's a long-running joke about no-one being able to remember Stata's "reshape" command. ([Exhibit A](#).)

It's easy to see this happening with the `pivot_*` functions too. However, I find that easier to remember the commands as long as I remember the argument order is *"names"* then *"values"*.

3) tidyr::separate

```
economists = data.frame(name = c("Adam.Smith", "Paul.Samuelson", "Milton.Friedman"))  
economists
```

```
##           name  
## 1   Adam.Smith  
## 2 Paul.Samuelson  
## 3 Milton.Friedman
```

```
economists > separate(name, c("first_name", "last_name"))
```

```
## first_name last_name  
## 1      Adam      Smith  
## 2      Paul Samuelson  
## 3      Milton  Friedman
```

This command is pretty smart. But to avoid ambiguity, you can also specify the separation character with `separate( ... , sep=".")`.

3) tidyr::separate cont.

A related function is `separate_rows`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data).

```
jobs = data.frame(
  name = c("Jack", "Jill"),
  occupation = c("Homemaker", "Philosopher, Philanthropist, Troublemaker")
)
jobs
```

```
##   name                occupation
## 1 Jack              Homemaker
## 2 Jill Philosopher, Philanthropist, Troublemaker
```

```
## Now split out Jill's various occupations into different rows
jobs > separate_rows(occupation)
```

```
## # A tibble: 4 × 2
##   name  occupation
##   <chr> <chr>
## 1 Jack  Homemaker
## 2 Jill  Philosopher
## 3 Jill  Philanthropist
## 4 Jill  Troublemaker
```

4) tidyr::unite

```
gdp = data.frame(  
  yr = rep(2016, times = 4),  
  mnth = rep(1, times = 4),  
  dy = 1:4,  
  gdp = rnorm(4, mean = 100, sd = 2)  
)  
gdp
```

```
##      yr mnth dy      gdp  
## 1 2016    1  1 98.99640  
## 2 2016    1  2 98.18651  
## 3 2016    1  3 100.61436  
## 4 2016    1  4 100.07875
```

```
## Combine "yr", "mnth", and "dy" into one "date" column  
gdp > unite(date, c("yr", "mnth", "dy"), sep = "-")
```

```
##      date      gdp  
## 1 2016-1-1 98.99640  
## 2 2016-1-2 98.18651  
## 3 2016-1-3 100.61436  
## 4 2016-1-4 100.07875
```


4) tidyr::unite cont.

Note that `unite` will automatically create a character variable. You can see this better if we convert it to a tibble.

```
gdp_u = gdp > unite(date, c("yr", "mnth", "dy"), sep = "-") > as_tibble()
gdp_u
```

```
## # A tibble: 4 × 2
##   date      gdp
##   <chr>    <dbl>
## 1 2016-1-1  99.0
## 2 2016-1-2  98.2
## 3 2016-1-3 101.
## 4 2016-1-4 100.
```

If you want to convert it to something else (e.g. date or numeric) then you will need to modify it using `mutate`. See the next slide for an example, using the `lubridate` package's super helpful date conversion functions.

4) tidyr::unite cont.

(continued from previous slide)

```
library(lubridate)
gdp_u > mutate(date = ymd(date))
```

```
## # A tibble: 4 × 2
##   date      gdp
##   <date>    <dbl>
## 1 2016-01-01  99.0
## 2 2016-01-02  98.2
## 3 2016-01-03 101.
## 4 2016-01-04 100.
```

Other tidyr goodies

Use `crossing` to get the full combination of a group of variables.¹

```
crossing(side=c("left", "right"), height=c("top", "bottom"))
```

```
## # A tibble: 4 × 2
##   side  height
##   <chr> <chr>
## 1 left  bottom
## 2 left  top
## 3 right bottom
## 4 right top
```

See `?expand` and `?complete` for more specialised functions that allow you to fill in (implicit) missing data or variable combinations in existing data frames.

- You'll encounter this during your next assignment.

¹ Base R alternative: `expand.grid`.

Summary

Key verbs

dplyr

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarize`

tidyr

1. `pivot_longer`
2. `pivot_wider`
3. `separate`
4. `unite`

Other useful items include: pipes (`>`), grouping (`group_by`), joining functions (`left_join`, `inner_join`, etc.).

Next lecture: Data cleaning and wrangling (2): data.table
