

# Data Science for Economists

## Lecture 9: Data cleaning & wrangling: Tidyverse

---

Drew Van Kuiken

University of North Carolina | ECON 370

# Table of contents

1. Prologue
2. Tidyverse basics
3. Data wrangling with dplyr
  - filter
  - arrange
  - select
  - mutate
  - summarize
  - joins
4. Data tidying with tidyr
  - pivot\_longer / pivot\_wider
  - separate and unite
5. Summary

# Prologue

---

# What is "tidy" data?

## Resources:

- [Vignette](#) (from the **tidyr** package)
- [Original paper](#) (Hadley Wickham, 2014 JSS)

## Key points:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Basically, tidy data is more likely to be [long \(i.e. narrow\) format](#) than wide format.

# Checklist

## R packages you'll need today

☑ **tidyverse**

☑ **nycflights13**

I'll hold off loading these libraries for now. But you can install/update them both with the following command.

```
install.packages(c('tidyverse', 'nycflights13'), repos = 'https://cran.rstudio.com', c
```

---

**Tip:** If you're on Linux, then I *strongly* recommend installing the pre-compiled binary versions of these packages from **RSPM** instead of CRAN. The exact repo mirror varies by distro (see the link). But on Ubuntu 20.04, for example, you'd use:

```
install.packages(c('tidyverse', 'nycflights13'), repos = 'https://packagemanager.rstu
```

---

# Tidyverse basics

---

# Tidyverse vs. base R

Much digital ink has been spilled over the "tidyverse vs. base R" debate.

I won't delve into this debate here, because I think the answer is **clear**: We should teach the tidyverse first (or, at least, early).

- The documentation and community support are outstanding.
- Having a consistent philosophy and syntax makes it easier to learn.
- Provides a convenient "front-end" to big data tools that we'll use later in the course.
- For data cleaning, wrangling, and plotting, the tidyverse really is a no-brainer.<sup>1</sup>

**But...** this certainly shouldn't put you off learning base R alternatives.

- Base R is extremely flexible and powerful (and stable).
- There are some things that you'll have to venture outside of the tidyverse for.
- A combination of tidyverse and base R is often the best solution to a problem.
- Excellent base R data manipulation tutorials: [here](#) and [here](#).

# Tidyverse vs. base R (cont.)

One point of convenience is that there is often a direct correspondence between a tidyverse command and its base R equivalent.

These generally follow a `tidyverse::snake_case` vs `base::period.case` rule. E.g. Compare:

tidyverse	base
<code>?readr::read_csv</code>	<code>?utils::read.csv</code>
<code>?dplyr::if_else</code>	<code>?base::ifelse</code>
<code>?tibble::tibble</code>	<code>?base::data.frame</code>

Etcetera.

If you call up the above examples, you'll see that the tidyverse alternative typically offers some enhancements or other useful options (and sometimes restrictions) over its base counterpart.

**Remember:** There are (almost) always multiple ways to achieve a single goal in R.



# Tidyverse packages

Let's load the tidyverse meta-package and check the output.

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats    1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2    3.5.1      ✓ tibble     3.2.1
## ✓ lubridate  1.9.3      ✓ tidyr      1.3.1
## ✓ purrr      1.0.2
## — Conflicts ————— tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to
```

We see that we have actually loaded a number of packages (which could also be loaded individually): **ggplot2**, **tibble**, **dplyr**, etc.

- We can also see information about the package versions and some namespace conflicts.

# Tidyverse packages (cont.)

The tidyverse actually comes with a lot more packages than those that are just loaded automatically.<sup>1</sup>

```
tidyverse_packages()
```

```
## [1] "broom"          "conflicted"    "cli"           "dbplyr"
## [5] "dplyr"          "dtplyr"        "forcats"       "ggplot2"
## [9] "googledrive"    "googlesheets4" "haven"         "hms"
## [13] "httr"           "jsonlite"      "lubridate"     "magrittr"
## [17] "modelr"         "pillar"        "purrr"         "ragg"
## [21] "readr"          "readxl"        "reprex"        "rlang"
## [25] "rstudioapi"     "rvest"         "stringr"       "tibble"
## [29] "tidyr"          "xml2"          "tidyverse"
```

- E.g. The **lubridate** package for working with dates and the **rvest** package for webscraping.
- However, bear in mind that these packages will have to be loaded separately.

<sup>1</sup> It also includes a *lot* of dependencies upon installation. This is a matter of some **controversy**.

# Tidyverse packages (cont.)

I hope to cover most of the tidyverse packages over the length of this course.

Today, however, I'm only really going to focus on two packages:

1. **dplyr**
2. **tidyr**

These are the workhorse packages for cleaning and wrangling data. They are thus the ones that you will likely make the most use of (alongside **ggplot2**, which we already met back in Lecture 1).

- Data cleaning and wrangling occupies an inordinate amount of time, no matter where you are in your research career.

# We're in (New) Pipe-land Now

The `tidyverse` uses a pipe `▷` (`|>`) which lets you send (i.e. "pipe") intermediate output to another command. Shortcut: **Ctrl/Command+Shift+M**

In other words, it allows us to chain together a sequence of simple operations and thereby implement a more complex operation.

I want to reiterate how cool pipes are, and how using them can dramatically improve the experience of reading and writing code. Compare:

```
## These next two lines of code do exactly the same thing.
```

```
mpg ▷ filter(manufacturer="audi") ▷ group_by(model) ▷ summarize(hwy_mean = mean(hwy))  
summarize(group_by(filter(mpg, manufacturer="audi"), model), hwy_mean = mean(hwy))
```

The first line reads from left to right, exactly how I thought of the operations in my head.

- Take this object (`mpg`), do this (`filter`), then do this (`group_by`), etc.

The second line totally inverts this logical order (the final operation comes first!)

- Who wants to read things inside out?

# An aside on pipes: |> (cont.)

The piped version of the code is even more readable if we write it over several lines. Here it is again and, this time, I'll run it for good measure so you can see the output:

```
mpg |>
  filter(manufacturer=="audi") |>
  group_by(model) |>
  summarize(hwy_mean = mean(hwy))
```

```
## # A tibble: 3 × 2
##   model      hwy_mean
##   <chr>      <dbl>
## 1 a4         28.3
## 2 a4 quattro 25.8
## 3 a6 quattro 24
```

Remember: Using vertical space costs nothing and makes for much more readable/writeable code than cramming things horizontally.

PS — The pipe is originally from the **magrittr** package, which can do some other cool things if you're inclined to explore. It's been so popular that even SQL is adopting a pipe syntax. SQL!

# dplyr

---

# Aside: dplyr 1.0.0 release

Some of the **dplyr** features that we'll cover today were introduced in **version 1.0.0** of the package.

- Version 1.0.0 is a big deal since it marks a stable code base for the package going forward.
- Please make sure that you are running at least **dplyr** 1.0.0 before continuing.

```
packageVersion('dplyr')
```

```
## [1] '1.1.4'
```

```
# install.packages('dplyr') ## install updated version if < 1.0.0
```

Note: **dplyr** 1.0.0 also notifies you about grouping variables every time you do operations on or with them. YMMV, but, personally, I find these messages annoying. You can **switch them off**.

```
options(dplyr.summarize.inform = FALSE) ## Add to .Rprofile to make permanent
```

# Key dplyr verbs

There are five key dplyr verbs that you need to learn.

1. `filter`: Filter (i.e. subset) rows based on their values.
2. `arrange`: Arrange (i.e. reorder) rows based on their values.
3. `select`: Select (i.e. subset) columns by their names:
4. `mutate`: Create new columns.
5. `summarize`: Collapse multiple rows into a single summary value.<sup>1</sup>

Let's practice these commands together using the `starwars` data frame that comes pre-packaged with dplyr.

<sup>1</sup> `summarise` with a "s" works too.



# 1) dplyr::filter

We can chain multiple filter commands with the pipe ( `>` ), or just separate them within a single filter command using commas.

```
starwars >
  filter(
    species = "Human",
    height ≥ 190
  )
```

```
## # A tibble: 4 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Darth Va...   202   136 none       white       yellow       41.9 male  masculi
## 2 Qui-Gon ...   193    89 brown      fair        blue         92  male  masculi
## 3 Dooku        193    80 white      fair        brown        102  male  masculi
## 4 Bail Pre...   191   NA black      tan         brown         67  male  masculi
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

What's the base R equivalent of this code?

# 1) dplyr::filter cont.

Regular expressions work well too.

```
starwars >
  filter(grepl("Skywalker", name))
```

```
## # A tibble: 3 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Luke Sky...   172    77 blond      fair       blue        19   male masculi...
## 2 Anakin S...   188    84 blond      fair       blue       41.9   male masculi...
## 3 Shmi Sky...   163    NA black      fair       brown       72   fema... femin...
```

## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## # vehicles <list>, starships <list>

# 1) dplyr::filter cont.

A very common `filter` use case is identifying (or removing) missing data cases.

```
starwars >
  filter(is.na(height))
```

```
## # A tibble: 6 × 14
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Arvel Cr...    NA    NA brown      fair      brown            NA male  mascu...
## 2 Finn          NA    NA black     dark     dark            NA male  mascu...
## 3 Rey           NA    NA brown     light    hazel            NA fema... femin...
## 4 Poe Dame...    NA    NA brown     light    brown            NA male  mascu...
## 5 BB8           NA    NA none      none     black            NA none  mascu...
## 6 Captain ...    NA    NA none      none     unknown          NA fema... femin...
## #   i 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

To remove missing observations, simply use negation: `filter(!is.na(height))`. Try this yourself.

## 2) dplyr::arrange

```
starwars ▶  
  arrange(birth_year)
```

```
## # A tibble: 87 × 14  
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender  
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  
## 1 Wicket ...      88   20  brown      brown      brown          8  male  masculi  
## 2 IG-88          200  140  none       metal      red           15  none  masculi  
## 3 Luke Sk...     172   77  blond      fair       blue          19  male  masculi  
## 4 Leia Or...     150   49  brown      light      brown          19  fema... feminin  
## 5 Wedge A...     170   77  brown      fair       hazel          21  male  masculi  
## 6 Plo Koon      188   80  none       orange     black          22  male  masculi  
## 7 Biggs D...     183   84  black      light      brown          24  male  masculi  
## 8 Han Solo       180   80  brown      fair       brown          29  male  masculi  
## 9 Lando C...     177   79  black      dark       brown          31  male  masculi  
## 10 Boba Fe...     183  78.2  black      fair       brown          31.5 male  masculi  
## # i 77 more rows  
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

Note: Arranging on a character-based column (i.e. strings) will sort alphabetically. Try this yourself by arranging according to the "name" column.

## 2) dplyr::arrange cont.

We can also arrange items in descending order using `arrange(desc())`.

```
starwars ▷
```

```
  arrange(desc(birth_year))
```

```
## # A tibble: 87 × 14
```

```
##   name      height  mass hair_color skin_color eye_color birth_year sex  gender
##   <chr>      <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>
## 1 Yoda         66    17 white      green      brown         896 male  masculi...
## 2 Jabba D...   175  1358 <NA>      green-tan... orange         600 herm... masculi...
## 3 Chewbac...   228   112 brown     unknown    blue          200 male  masculi...
## 4 C-3PO        167    75 <NA>      gold       yellow         112 none  masculi...
## 5 Dooku         193    80 white     fair       brown         102 male  masculi...
## 6 Qui-Gon...   193    89 brown     fair       blue           92 male  masculi...
## 7 Ki-Adi-...   198    82 white     pale       yellow          92 male  masculi...
## 8 Finis V...   170    NA blond     fair       blue           91 male  masculi...
## 9 Palpati...   170    75 grey      pale       yellow          82 male  masculi...
## 10 Cliegg ...   183    NA brown     fair       blue           82 male  masculi...
## # i 77 more rows
## # i 5 more variables: homeworld <chr>, species <chr>, films <list>,
## #   vehicles <list>, starships <list>
```

### 3) dplyr::select

Use commas to select multiple columns out of a data frame. (You can also use "first:last" for consecutive columns). Deselect a column with "-".

```
starwars ►  
  select(name:skin_color, species, -height)
```

```
## # A tibble: 87 × 5  
##   name                mass hair_color  skin_color species  
##   <chr>              <dbl> <chr>      <chr>      <chr>  
## 1 Luke Skywalker      77 blond      fair        Human  
## 2 C-3PO                75 <NA>      gold        Droid  
## 3 R2-D2                32 <NA>      white, blue Droid  
## 4 Darth Vader         136 none      white        Human  
## 5 Leia Organa          49 brown      light        Human  
## 6 Owen Lars           120 brown, grey light        Human  
## 7 Beru Whitesun Lars   75 brown      light        Human  
## 8 R5-D4                32 <NA>      white, red   Droid  
## 9 Biggs Darklighter   84 black      light        Human  
## 10 Obi-Wan Kenobi      77 auburn, white fair         Human  
## # i 77 more rows
```

### 3) dplyr::select *cont.*

You can also rename some (or all) of your selected variables in place.

```
starwars ►  
  select(alias=name, crib=homeworld, sex=gender)
```

```
## # A tibble: 87 × 3  
##   alias          crib      sex  
##   <chr>         <chr>   <chr>  
## 1 Luke Skywalker Tatooine masculine  
## 2 C-3PO         Tatooine masculine  
## 3 R2-D2         Naboo    masculine  
## 4 Darth Vader   Tatooine masculine  
## 5 Leia Organa   Alderaan feminine  
## 6 Owen Lars     Tatooine masculine  
## 7 Beru Whitesun Lars Tatooine feminine  
## 8 R5-D4         Tatooine masculine  
## 9 Biggs Darklighter Tatooine masculine  
## 10 Obi-Wan Kenobi Stewjon  masculine  
## # i 77 more rows
```

If you just want to rename columns without subsetting them, you can use `rename`. Try this now by replacing `select( ... )` in the above code chunk with `rename( ... )`.

### 3) dplyr::select cont.

The `select(contains(PATTERN))` option provides a nice shortcut in relevant cases.

```
starwars ▷
```

```
  select(name, contains("color"))
```

```
## # A tibble: 87 × 4
```

##	name	hair_color	skin_color	eye_color
##	<chr>	<chr>	<chr>	<chr>
## 1	Luke Skywalker	blond	fair	blue
## 2	C-3PO	<NA>	gold	yellow
## 3	R2-D2	<NA>	white, blue	red
## 4	Darth Vader	none	white	yellow
## 5	Leia Organa	brown	light	brown
## 6	Owen Lars	brown, grey	light	blue
## 7	Beru Whitesun Lars	brown	light	blue
## 8	R5-D4	<NA>	white, red	red
## 9	Biggs Darklighter	black	light	brown
## 10	Obi-Wan Kenobi	auburn, white	fair	blue-gray
## #	i 77 more rows			



### 3) dplyr::select *cont.*

The `select(..., everything())` option is another useful shortcut if you only want to bring some variable(s) to the "front" of a data frame.

```
starwars ►  
  select(species, homeworld, everything()) ►  
  head(5)
```

```
## # A tibble: 5 × 14  
##   species homeworld name          height  mass hair_color skin_color eye_color  
##   <chr>    <chr>    <chr>          <int> <dbl> <chr>      <chr>    <chr>  
## 1 Human   Tatooine  Luke Skywalker    172    77 blond      fair      blue  
## 2 Droid   Tatooine  C-3PO             167    75 <NA>      gold      yellow  
## 3 Droid   Naboo     R2-D2              96    32 <NA>      white, blue red  
## 4 Human   Tatooine  Darth Vader       202   136 none       white      yellow  
## 5 Human   Alderaan  Leia Organa       150    49 brown      light      brown  
## # i 6 more variables: birth_year <dbl>, sex <chr>, gender <chr>, films <list>,  
## #   vehicles <list>, starships <list>
```

Note: The `relocate` function available in dplyr 1.0.0 has brought a lot more functionality to ordering of columns. See [here](#).

## 4) dplyr::mutate

You can create new columns from scratch, or (more commonly) as transformations of existing columns.

```
starwars >
  select(name, birth_year) >
  mutate(dog_years = birth_year * 7) >
  mutate(comment = paste0(name, " is ", dog_years, " in dog years."))
```

```
## # A tibble: 87 × 4
##   name          birth_year dog_years comment
##   <chr>          <dbl>     <dbl> <chr>
## 1 Luke Skywalker      19       133 Luke Skywalker is 133 in dog years.
## 2 C-3P0             112       784 C-3P0 is 784 in dog years.
## 3 R2-D2              33       231 R2-D2 is 231 in dog years.
## 4 Darth Vader       41.9      293.3 Darth Vader is 293.3 in dog years.
## 5 Leia Organa        19       133 Leia Organa is 133 in dog years.
## 6 Owen Lars          52       364 Owen Lars is 364 in dog years.
## 7 Beru Whitesun Lars  47       329 Beru Whitesun Lars is 329 in dog yea...
## 8 R5-D4              NA         NA R5-D4 is NA in dog years.
## 9 Biggs Darklighter  24       168 Biggs Darklighter is 168 in dog year...
## 10 Obi-Wan Kenobi     57       399 Obi-Wan Kenobi is 399 in dog years.
## # i 77 more rows
```

## 4) dplyr::mutate cont.

Note: `mutate` is order aware. So you can chain multiple mutates in a single call.

```
starwars >
  select(name, birth_year) >
  mutate(
    dog_years = birth_year * 7, ## Separate with a comma
    comment = paste0(name, " is ", dog_years, " in dog years.")
  )
```

```
## # A tibble: 87 × 4
```

	name	birth_year	dog_years	comment
	<chr>	<dbl>	<dbl>	<chr>
## 1	Luke Skywalker	19	133	Luke Skywalker is 133 in dog years.
## 2	C-3PO	112	784	C-3PO is 784 in dog years.
## 3	R2-D2	33	231	R2-D2 is 231 in dog years.
## 4	Darth Vader	41.9	293.	Darth Vader is 293.3 in dog years.
## 5	Leia Organa	19	133	Leia Organa is 133 in dog years.
## 6	Owen Lars	52	364	Owen Lars is 364 in dog years.
## 7	Beru Whitesun Lars	47	329	Beru Whitesun Lars is 329 in dog yea...
## 8	R5-D4	NA	NA	R5-D4 is NA in dog years.
## 9	Biggs Darklighter	24	168	Biggs Darklighter is 168 in dog year...
## 10	Obi-Wan Kenobi	57	399	Obi-Wan Kenobi is 399 in dog years.
## #	i 77 more rows			

## 4) dplyr::mutate cont.

Boolean, logical and conditional operators all work well with `mutate` too.

```
starwars >
  select(name, height) >
  filter(name %in% c("Luke Skywalker", "Anakin Skywalker")) >
  mutate(tall1 = height > 180) >
  mutate(tall2 = ifelse(height > 180, "Tall", "Short")) ## Same effect, but can choose
```

```
## # A tibble: 2 × 4
##   name          height tall1 tall2
##   <chr>         <int> <lgl> <chr>
## 1 Luke Skywalker    172 FALSE Short
## 2 Anakin Skywalker    188  TRUE  Tall
```

## 4) dplyr::mutate *cont.*

Lastly, combining `mutate` with the `across` feature in dplyr 1.0.0+ allows you to easily work on a subset of variables. For example:

```
starwars >
  select(name:eye_color) >
  mutate(across(where(is.character), toupper)) >
  head(5)
```

```
## # A tibble: 5 × 6
##   name          height  mass hair_color skin_color eye_color
##   <chr>         <int> <dbl> <chr>      <chr>      <chr>
## 1 LUKE SKYWALKER   172    77 BLOND      FAIR        BLUE
## 2 C-3PO           167    75 <NA>      GOLD        YELLOW
## 3 R2-D2            96    32 <NA>      WHITE, BLUE RED
## 4 DARTH VADER     202   136 NONE      WHITE        YELLOW
## 5 LEIA ORGANA     150    49 BROWN     LIGHT        BROWN
```

*Note:* This workflow (i.e. combining `mutate` and `across`) supersedes the old "scoped" variants of `mutate` that you might have used previously. More details [here](#) and [here](#).

## 5) dplyr::summarize

Particularly useful in combination with the `group_by` command.

```
starwars >
  group_by(species, gender) >
  summarize(mean_height = mean(height, na.rm = TRUE))
```

## `summarise()` has grouped output by 'species'. You can override using the  
## `.groups` argument.

```
## # A tibble: 42 × 3
## # Groups:   species [38]
##   species    gender  mean_height
##   <chr>      <chr>      <dbl>
## 1 Aleena     masculine      79
## 2 Besalisk   masculine     198
## 3 Cerean     masculine     198
## 4 Chagrian   masculine     196
## 5 Clawdite   feminine     168
## 6 Droid      feminine      96
## 7 Droid      masculine     140
## 8 Dug        masculine     112
## 9 Ewok       masculine      88
## 10 Geonosian masculine     183
## # i 32 more rows
```

## 5) dplyr::summarize *cont.*

Note that including "na.rm = TRUE" (or, its alias "na.rm = T") is usually a good idea with summarize functions. Otherwise, any missing value will propagate to the summarized value too.

```
## Probably not what we want
```

```
starwars %>
```

```
  summarize(mean_height = mean(height))
```

```
## # A tibble: 1 × 1
```

```
##   mean_height
```

```
##         <dbl>
```

```
## 1          NA
```

```
## Much better
```

```
starwars %>
```

```
  summarize(mean_height = mean(height, na.rm = TRUE))
```

```
## # A tibble: 1 × 1
```

```
##   mean_height
```

```
##         <dbl>
```

```
## 1        175.
```

## 5) dplyr::summarize *cont.*

The same `across`-based workflow that we saw with `mutate` a few slides back also works with `summarize`. For example:

```
starwars >
  group_by(species) >
  summarize(across(where(is.numeric), function(x) mean(x, na.rm=T))) >
  head(5)
```

```
## # A tibble: 5 × 4
##   species height  mass birth_year
##   <chr>    <dbl> <dbl>      <dbl>
## 1 Aleena      79     15         NaN
## 2 Besalisk   198    102         NaN
## 3 Cerean     198     82          92
## 4 Chagrian   196    NaN         NaN
## 5 Clawdite   168     55         NaN
```

*Note:* Again, this functionality supersedes the old "scoped" variants of `summarize` that you used prior to dplyr 1.0.0. Details [here](#) and [here](#).



# Other dplyr goodies

`group_by` and `ungroup`: For (un)grouping.

- Particularly useful with the `summarize` and `mutate` commands, as we've already seen.

`slice`: Subset rows by position rather than filtering by values.

- E.g. `starwars > slice(c(1, 5))`

`pull`: Extract a column from a data frame as a vector or scalar.

- E.g. `starwars > filter(gender="female") > pull(height)`

`count` and `distinct`: Number and isolate unique observations.

- E.g. `starwars > count(species)`, or `starwars > distinct(species)`
- You could also use a combination of `mutate`, `group_by`, and `n()`, e.g. `starwars > group_by(species) > mutate(num = n())`.

# Other dplyr goodies (cont.)

There are also a whole class of **window functions** for getting leads and lags, ranking, creating cumulative aggregates, etc.

- See `vignette("window-functions")`.

The final set of dplyr "goodies" are the family of join operations. However, these are important enough that I want to go over some concepts in a bit more depth...


- We will encounter and practice these many more times as the course progresses.

# Joins

One of the mainstays of the dplyr package is merging data with the family **join operations**.

- `inner_join(df1, df2)`
- `left_join(df1, df2)`
- `right_join(df1, df2)`
- `full_join(df1, df2)`
- `semi_join(df1, df2)`
- `anti_join(df1, df2)`

# Visualizing Joins



For the simple examples that I'm going to show here, we'll need some data sets that come bundled with the **nycflights13** package.

- Load it now and then inspect these data frames in your own console.

```
library(nycflights13)
flights
planes
```

# Joins (cont.)

Let's perform a **left join** on the flights and planes datasets.

- *Note:* I'm going subset columns after the join, but only to keep text on the slide.

```
left_join(flights, planes) ▷  
  select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, type, model)
```

```
## Joining with by = join_by(year, tailnum)
```

```
## # A tibble: 336,776 × 10
```

```
##   year month   day dep_time arr_time carrier flight tailnum type  model  
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>   <chr> <chr>  
## 1  2013     1     1     517     830 UA      1545 N14228 <NA> <NA>  
## 2  2013     1     1     533     850 UA      1714 N24211 <NA> <NA>  
## 3  2013     1     1     542     923 AA      1141 N619AA <NA> <NA>  
## 4  2013     1     1     544    1004 B6       725 N804JB <NA> <NA>  
## 5  2013     1     1     554     812 DL       461 N668DN <NA> <NA>  
## 6  2013     1     1     554     740 UA      1696 N39463 <NA> <NA>  
## 7  2013     1     1     555     913 B6       507 N516JB <NA> <NA>  
## 8  2013     1     1     557     709 EV      5708 N829AS <NA> <NA>  
## 9  2013     1     1     557     838 B6        79 N593JB <NA> <NA>  
## 10 2013     1     1     558     753 AA       301 N3ALAA <NA> <NA>
```

```
## # i 336,766 more rows
```

# Joins (cont.)

*(continued from previous slide)*

Note that dplyr made a reasonable guess about which columns to join on (i.e. columns that share the same name). It also told us its choices:

```
## Joining, by = c("year", "tailnum")
```

However, there's an obvious problem here: the variable "year" does not have a consistent meaning across our joining datasets!

- In one it refers to the *year of flight*, in the other it refers to *year of construction*.

Luckily, there's an easy way to avoid this problem.

- See if you can figure it out before turning to the next slide.
- Try `?dplyr::join`.

# Joins (cont.)

(continued from previous slide)

You just need to be more explicit in your join call by using the `by =` argument.

- You can also rename any ambiguous columns to avoid confusion.

```
left_join(
  flights,
  planes ▷ rename(year_built = year), ## Not necessary w/ below line, but helpful
  by = "tailnum" ## Be specific about the joining column
) ▷
select(year, month, day, dep_time, arr_time, carrier, flight, tailnum, year_built, 1
head(3) ## Just to save vertical space on the slide
```

```
## # A tibble: 3 × 11
##   year month   day dep_time arr_time carrier flight tailnum year_built type
##   <int> <int> <int>   <int>   <int> <chr>   <int> <chr>      <int> <chr>
## 1  2013     1     1     517     830 UA      1545 N14228     1999 Fixed w...
## 2  2013     1     1     533     850 UA      1714 N24211     1998 Fixed w...
## 3  2013     1     1     542     923 AA      1141 N619AA     1990 Fixed w...
## # i 1 more variable: model <chr>
```

# Joins (cont.)

(continued from previous slide)

Last thing I'll mention for now; note what happens if we again specify the join column... but don't rename the ambiguous "year" column in at least one of the given data frames.

```
left_join(
  flights,
  planes, ## Not renaming "year" to "year_built" this time
  by = "tailnum"
) ▷
select(contains("year"), month, day, dep_time, arr_time, carrier, flight, tailnum, 1
head(3)
```

```
## # A tibble: 3 × 11
##   year.x year.y month   day dep_time arr_time carrier flight tailnum type  model
##   <int> <int> <int> <int>   <int>   <int> <chr>   <int> <chr>   <chr> <chr>
## 1  2013  1999     1     1     517     830 UA       1545 N14228 Fixe... 737-...
## 2  2013  1998     1     1     533     850 UA       1714 N24211 Fixe... 737-...
## 3  2013  1990     1     1     542     923 AA       1141 N619AA Fixe... 757-...
```

Make sure you know what "year.x" and "year.y" are. Again, it pays to be specific.



# tidyr

---

# Key tidyr verbs

1. `pivot_longer`: Pivot wide data into long format (i.e. "melt").<sup>1</sup>
2. `pivot_wider`: Pivot long data into wide format (i.e. "cast").<sup>2</sup>
3. `separate`: Separate (i.e. split) one column into multiple columns.
4. `unite`: Unite (i.e. combine) multiple columns into one.

Let's practice these verbs together in class.

- Side question: Which of `pivot_longer` vs `pivot_wider` produces "tidy" data?

<sup>1</sup> Updated version of `tidyr::gather`.

<sup>2</sup> Updated version of `tidyr::spread`.

# 1) tidyr::pivot\_longer

```
stocks = data.frame( ## Could use "tibble" instead of "data.frame" if you prefer
  time = as.Date('2009-01-01') + 0:1,
  X = rnorm(2, 0, 1),
  Y = rnorm(2, 0, 2),
  Z = rnorm(2, 0, 4)
)
stocks
```

```
##           time           X           Y           Z
## 1 2009-01-01 1.944249 -1.613529 -7.097121
## 2 2009-01-02 1.561860  1.395649  2.988417
```

```
stocks > pivot_longer(-time, names_to="stock", values_to="price")
```

```
## # A tibble: 6 × 3
##   time           stock price
##   <date>         <chr> <dbl>
## 1 2009-01-01 X         1.94
## 2 2009-01-01 Y        -1.61
## 3 2009-01-01 Z        -7.10
## 4 2009-01-02 X         1.56
## 5 2009-01-02 Y         1.40
## 6 2009-01-02 Z         2.99
```

# 1) tidyr::pivot\_longer cont.

Let's quickly save the "tidy" (i.e. long) stocks data frame for use on the next slide.

```
## Write out the argument names this time: i.e. "names_to=" and "values_to="  
tidy_stocks =  
  stocks ▷  
  pivot_longer(-time, names_to="stock", values_to="price")
```

## 2) tidyr::pivot\_wider

```
tidy_stocks > pivot_wider(names_from=stock, values_from=price)
```

```
## # A tibble: 2 × 4
##   time          X      Y      Z
##   <date>      <dbl> <dbl> <dbl>
## 1 2009-01-01  1.94 -1.61 -7.10
## 2 2009-01-02  1.56  1.40  2.99
```

```
tidy_stocks > pivot_wider(names_from=time, values_from=price)
```

```
## # A tibble: 3 × 3
##   stock 2009-01-01 2009-01-02
##   <chr>      <dbl>      <dbl>
## 1 X          1.94          1.56
## 2 Y          -1.61          1.40
## 3 Z          -7.10          2.99
```

Note that the second example — which has combined different pivoting arguments — has effectively transposed the data.

# Aside: Remembering the `pivot_*` syntax

There's a long-running joke about no-one being able to remember Stata's "reshape" command. ([Exhibit A](#).)

It's easy to see this happening with the `pivot_*` functions too. However, I find that easier to remember the commands as long as I remember the argument order is *"names"* then *"values"*.

### 3) tidyr::separate

```
economists = data.frame(name = c("Adam.Smith", "Paul.Samuelson", "Milton.Friedman"))  
economists
```

```
##           name  
## 1   Adam.Smith  
## 2 Paul.Samuelson  
## 3 Milton.Friedman
```

```
economists > separate(name, c("first_name", "last_name"))
```

```
## first_name last_name  
## 1      Adam      Smith  
## 2      Paul Samuelson  
## 3      Milton  Friedman
```

This command is pretty smart. But to avoid ambiguity, you can also specify the separation character with `separate( ... , sep=".")`.

### 3) tidyr::separate cont.

A related function is `separate_rows`, for splitting up cells that contain multiple fields or observations (a frustratingly common occurrence with survey data).

```
jobs = data.frame(
  name = c("Jack", "Jill"),
  occupation = c("Homemaker", "Philosopher, Philanthropist, Troublemaker")
)
jobs
```

```
##   name                occupation
## 1 Jack                Homemaker
## 2 Jill Philosopher, Philanthropist, Troublemaker
```

```
## Now split out Jill's various occupations into different rows
jobs > separate_rows(occupation)
```

```
## # A tibble: 4 × 2
##   name  occupation
##   <chr> <chr>
## 1 Jack  Homemaker
## 2 Jill  Philosopher
## 3 Jill  Philanthropist
## 4 Jill  Troublemaker
```



## 4) tidyr::unite

```
gdp = data.frame(  
  yr = rep(2016, times = 4),  
  mnth = rep(1, times = 4),  
  dy = 1:4,  
  gdp = rnorm(4, mean = 100, sd = 2)  
)  
gdp
```

```
##      yr mnth dy      gdp  
## 1 2016     1  1 99.28177  
## 2 2016     1  2 96.99765  
## 3 2016     1  3 99.16636  
## 4 2016     1  4 97.02445
```

```
## Combine "yr", "mnth", and "dy" into one "date" column  
gdp > unite(date, c("yr", "mnth", "dy"), sep = "-")
```

```
##      date      gdp  
## 1 2016-1-1 99.28177  
## 2 2016-1-2 96.99765  
## 3 2016-1-3 99.16636  
## 4 2016-1-4 97.02445
```

## 4) tidyr::unite cont.

Note that `unite` will automatically create a character variable. You can see this better if we convert it to a tibble.

```
gdp_u = gdp > unite(date, c("yr", "mnth", "dy"), sep = "-") > as_tibble()
gdp_u
```

```
## # A tibble: 4 × 2
##   date      gdp
##   <chr>    <dbl>
## 1 2016-1-1  99.3
## 2 2016-1-2  97.0
## 3 2016-1-3  99.2
## 4 2016-1-4  97.0
```

If you want to convert it to something else (e.g. date or numeric) then you will need to modify it using `mutate`. See the next slide for an example, using the `lubridate` package's super helpful date conversion functions.

## 4) tidyr::unite cont.

*(continued from previous slide)*

```
library(lubridate)
gdp_u > mutate(date = ymd(date))
```

```
## # A tibble: 4 × 2
##   date      gdp
##   <date>    <dbl>
## 1 2016-01-01  99.3
## 2 2016-01-02  97.0
## 3 2016-01-03  99.2
## 4 2016-01-04  97.0
```

# Other tidyr goodies

Use `crossing` to get the full combination of a group of variables.<sup>1</sup>

```
crossing(side=c("left", "right"), height=c("top", "bottom"))
```

```
## # A tibble: 4 × 2
##   side  height
##   <chr> <chr>
## 1 left  bottom
## 2 left  top
## 3 right bottom
## 4 right top
```

See `?expand` and `?complete` for more specialised functions that allow you to fill in (implicit) missing data or variable combinations in existing data frames.

- You'll encounter this during your next assignment.

<sup>1</sup> Base R alternative: `expand.grid`.

# Summary

---

# Key verbs

## dplyr

1. `filter`
2. `arrange`
3. `select`
4. `mutate`
5. `summarize`

## tidyr

1. `pivot_longer`
2. `pivot_wider`
3. `separate`
4. `unite`

Other useful items include: pipes ( `>` ), grouping ( `group_by` ), joining functions ( `left_join`, `inner_join`, etc.).

Next lecture: Data cleaning and  
wrangling: (2) data.table

---