

# Data Science for Economists

## Lecture 11: Webscraping part 1 - CSS

---

Drew Van Kuiken

University of North Carolina | ECON 370

# The Internet as Data

Let's get it onto our computers.

For today, load **rvest** and **janitor** into your R session, alongside **tidyverse**, **lubridate**, **data.table**, and **hrbrthemes**

```
## Load and install the packages that we'll be using today  
if (!require("pacman")) install.packages("pacman")  
pacman::p_load(tidyverse, rvest, lubridate, janitor, data.table, hrbrthemes)  
## ggplot2 plotting theme (optional)  
theme_set(hrbrthemes::theme_ipsum())
```

Much credit to Grant McDermott for the content on these slides.

# Internet Data Is Stored in 2 Ways:

## 1. Server-side

- Data stored at the server, which sends HTML code with data in it to us
- **Our process:** trudging through CSS selectors

## 2. Client-side

- Our browser requests data from the server, server sends specific info we asked for
- **Our process:** pinging an API endpoint

This lecture will focus on server-side scraping; we'll do client-side scraping next

# Before we get started

1. Be a good internet user
2. It's easy to accidentally kill some poor website
3. It's probably legal?

Our main package today is **rvest**, part of the tidyverse. Based on **Beautiful Soup**

# HTML Basics

Here's some simple HTML:

```
<html>
<head>
  <title>Page title</title>
</head>
<body>
  <h1 id='first'>A heading</h1>
  <p>Some text &amp; <b>some bold text.</b></p>
  <img src='myimg.png' width='100' height='100'>
</body>
```

# HTML Basics

Here's some simple HTML:

```
<html>
<head> ## head is an element #<<
  <title>Page title</title>
</head>
<body>
  <h1 id='first'>A heading</h1>
  <p>Some text &amp; <b>some bold text.</b></p>
  <img src='myimg.png' width='100' height='100'>
</body>
```

# HTML Basics

Here's some simple HTML:

```
<html>
<head>
  <title>Page title</title>
</head>
<body> ## as is body #<<
  <h1 id='first'>A heading</h1>
  <p>Some text &amp; <b>some bold text.</b></p>
  <img src='myimg.png' width='100' height='100'>
</body>
```

# HTML Basics

Here's some simple HTML:

```
<html>
<head>
  <title>Page title</title>
</head>
<body>
  <h1 id='first'>A heading</h1> ## id='first' is an attribute #<<
  <p>Some text &amp; <b>some bold text.</b></p>
  <img src='myimg.png' width='100' height='100'>
</body>
```



# HTML Basics

Here's some simple HTML:

```
<html>
<head>
  <title>Page title</title>
</head>
<body>
  <h1 id='first'>A heading</h1>
  <p>Some text &amp; <b>some bold text.</b></p> ## stuff in between is contents #<<
  <img src='myimg.png' width='100' height='100'>
</body>
```

Lucky for us: we don't need to write HTML. Just read it.

# HTML Basics - Tags

- Tags all start with `<tag>` and end with `</tag>`
- Every page is in an `<html>` element with 2 children:
  - `<head>` contains metadata
  - `<body>` contains content you see
- Block tags form the overall structure of a page
  - `<h1>` provides a heading
  - `<p>` is a paragraph
- Inline tags like `<b>` (bold), `<i>` (italics), and `<a>` (links) exist
- Just a sample of tags, can look up others you don't know
- The rest is the content

# Example 1: Scraping Wikipedia

Let's imagine we want to scrape the [Men's 100 metres world record progression](https://en.wikipedia.org/wiki/Men%27s_100_metres_world_record_progression) page on Wikipedia.

In particular, we want to get the information from the 3 main tables on the webpage.

Here's what happens when we give R no instructions at all:

```
raw_wiki ← read_html("https://en.wikipedia.org/wiki/Men%27s_100_metres_world_record_100m")
raw_wiki
```

```
## {html_document}
## <html class="client-nojs vector-feature-language-in-header-enabled vector-feature-language-in-header-global" ...
## [1] <head>\n<meta http-equiv="Content-Type" content="text/html; charset=UTF-8 ...
## [2] <body class="skin--responsive skin-vector skin-vector-search-vue mediawik ...
```

```
class(raw_wiki)
```

```
## [1] "xml_document" "xml_node"
```

# Parsing HTML: Inspecting Web Pages

We can use `inspect` to get a better sense of what is available on a given webpage. What tags can we use to grab the desired information from our wikipedia page?

It looks like we want to grab information from tables with the class "wikitable."

# Brief Aside: CSS Selectors

We can parse HTML using **CSS Selectors**, which define patterns for locating HTML elements.

CSS Selectors are pretty complex, and we're going to keep it light today. If you want to learn more, check out the [CSS Diner](#). Additionally, if CSS Selectors are giving you tons of trouble, check out [SelectorGadget](#).

4 selectors to know:

1. `p` selects all `<p>` elements
2. `.title` selects all elements with `class` "title"
3. `p.special` selects all `<p>` elements with `class` "special"
4. `#title` selects the unique element with id attribute that equals "title"

# Getting Our Tables

We can use the selector `table` to select all `<table>` elements:

```
raw_wiki > html_elements("table")

## {xml_nodeset (15)}
## [1] <table class="box-Unreferenced_section plainlinks metadata ambox ambox-c ...
## [2] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Athlete\n</ ...
## [3] <table class="wikitable" style="text-align: left;"><tbody>\n<tr>\n<td st ...
## [4] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Wind\n</th> ...
## [5] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Wind\n</th> ...
## [6] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Athlete\n</ ...
## [7] <table class="nowraplinks mw-collapsible autocollapse navbox-inner" styl ...
## [8] <table class="nowraplinks navbox-subgroup" style="border-spacing:0"><tbo ...
## [9] <table class="nowraplinks hlist mw-collapsible autocollapse navbox-inner ...
## [10] <table class="nowraplinks navbox-subgroup" style="border-spacing:0"><tbo ...
## [11] <table class="nowraplinks navbox-subgroup" style="border-spacing:0"><tbo ...
## [12] <table class="nowraplinks mw-collapsible uncollapsed navbox-inner" style ...
## [13] <table class="nowraplinks navbox-subgroup" style="border-spacing:0"><tbo ...
## [14] <table class="nowraplinks navbox-subgroup" style="border-spacing:0"><tbo ...
## [15] <table class="nowraplinks navbox-subgroup" style="border-spacing:0"><tbo ...
```

Alas, so are many other things.

# Trying again

We want tables with class wikitable: `table.wikitable` should work!

```
raw_wiki ▷ html_elements("table.wikitable")
```

```
## {xml_nodeset (5)}  
## [1] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Athlete\n</th> ...  
## [2] <table class="wikitable" style="text-align: left;"><tbody>\n<tr>\n<td sty ...  
## [3] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Wind\n</th>\n ...  
## [4] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Wind\n</th>\n ...  
## [5] <table class="wikitable"><tbody>\n<tr>\n<th>Time\n</th>\n<th>Athlete\n</th> ...
```

Looks better!

# Accessing the actual info

We still don't have the table info. `html_table()` can help:

```
table_dfs ← raw_wiki ▷ html_elements("table.wikitable") ▷  
  html_table()  
table_dfs[[1]]
```

```
## # A tibble: 21 × 5  
##   Time Athlete      Nationality `Location of races`   Date  
##   <dbl> <chr>      <chr>      <chr>      <chr>  
## 1  10.8 Luther Cary    United States  Paris, France    July 4, 1...  
## 2  10.8 Cecil Lee      United Kingdom Brussels, Belgium September...  
## 3  10.8 Étienne De Ré   Belgium       Brussels, Belgium August 4,...  
## 4  10.8 L. Atcherley    United Kingdom Frankfurt/Main, Germany April 13,...  
## 5  10.8 Harry Beaton    United Kingdom Rotterdam, Netherlands August 28...  
## 6  10.8 Harald Anderson-Arbin Sweden         Helsingborg, Sweden August 9,...  
## 7  10.8 Isaac Westergren Sweden         Gävle, Sweden    September...  
## 8  10.8 Isaac Westergren Sweden         Gävle, Sweden    September...  
## 9  10.8 Frank Jarvis    United States  Paris, France    July 14, ...  
## 10 10.8 Walter Tewksbury United States  Paris, France    July 14, ...  
## # i 11 more rows
```

...this is crazy! `R` is cool sometimes.



# General Workflow

Your workflow using `rvest`: get html --> get desired html elements --> break into individual elements --> turn into table/text/numbers/whatever.

Useful commands for getting individual elements:

- `html_text2()` gets plain text contents of an HTML element
- `html_attr()` gets attributes from an HTML element (e.g., links)
- `html_table()` creates a data.frame from a table in an HTML element

# A Little Cleanup

Let's get our data frames in working order here:

```
table_dfs_int <- raw_wiki %>% html_elements("table.wikitable") %>%  
  html_table()
```

```
table_dfs <- lapply(table_dfs_int[c(1,3,4)], # drop unwanted tables  
  function(x) x %>%
```

```
    clean_names() %>% ## fix colnames, from the janitor package  
    mutate(date = mdy(date))) ## from lubridate
```

```
table_dfs[[1]]
```

```
## # A tibble: 21 × 5
```

##	time	athlete	nationality	location_of_races	date
##	<dbl>	<chr>	<chr>	<chr>	<date>
## 1	10.8	Luther Cary	United States	Paris, France	1891-07-04
## 2	10.8	Cecil Lee	United Kingdom	Brussels, Belgium	1892-09-25
## 3	10.8	Étienne De Ré	Belgium	Brussels, Belgium	1893-08-04
## 4	10.8	L. Atcherley	United Kingdom	Frankfurt/Main, Germany	1895-04-13
## 5	10.8	Harry Beaton	United Kingdom	Rotterdam, Netherlands	1895-08-28
## 6	10.8	Harald Anderson-Arbin	Sweden	Helsingborg, Sweden	1896-08-09
## 7	10.8	Isaac Westergren	Sweden	Gävle, Sweden	1898-09-11
## 8	10.8	Isaac Westergren	Sweden	Gävle, Sweden	1899-09-10
## 9	10.8	Frank Jarvis	United States	Paris, France	1900-07-14
## 10	10.8	Walter Tewksbury	United States	Paris, France	1900-07-14


# Combined Data

```
wr100 <- rbind(
  table_dfs[[1]] > select(time, athlete, nationality, date) >
    mutate(era="Pre-IAAF"),
  table_dfs[[2]] > select(time, athlete, nationality, date) >
    mutate(era="Pre-automatic"),
  table_dfs[[3]] > select(time, athlete, nationality, date) >
    mutate(era="Modern")
)
head(wr100)
```

```
## # A tibble: 6 × 5
```

##	time	athlete	nationality	date	era
##	<dbl>	<chr>	<chr>	<date>	<chr>
## 1	10.8	Luther Cary	United States	1891-07-04	Pre-IAAF
## 2	10.8	Cecil Lee	United Kingdom	1892-09-25	Pre-IAAF
## 3	10.8	Étienne De Ré	Belgium	1893-08-04	Pre-IAAF
## 4	10.8	L. Atcherley	United Kingdom	1895-04-13	Pre-IAAF
## 5	10.8	Harry Beaton	United Kingdom	1895-08-28	Pre-IAAF
## 6	10.8	Harald Anderson-Arbin	Sweden	1896-08-09	Pre-IAAF

# Dot Plot



That was easy. How about something harder.

# Craigslist - Watch Prices

Let's take a look at watch prices on [Craigslist in Raleigh](#).

We want to know: listing names, prices, locations, and links for the Raleigh craigslist.

Use the shell code below to poke around and see if you can download the information we want.

```
# read in html, get listing info
web ← "https://raleigh.craigslist.org/search/jwa?query=watch#search=1~gallery~0~0"
craigslist_listings ← read_html(web)▷
  html_elements("[INSERT ELEMENTS TO GET CODE HERE]")
craigslist_listings[[1]]
```

(Remember our basic strategy: `read_html` --> `html_elements` --> `html_element` --> `html_text2`. We're covering steps 1 and 2 here, don't worry about having your output look perfect.)

# My Solution

(Admittedly, it took a while for me to figure this out. Art, not science.)

All kinds of info is saved in `li a`!

We can use `html_elements()` to find an element that corresponds to each listing, then use `html_element()` to extract each individual variable:

```
# read in html, get listing info
web ← "https://raleigh.craigslist.org/search/jwa?query=watch#search=1~gallery~0~0"
craigslist_listings ← read_html(web) ▷ html_elements("li a")
craigslist_listings[[1]]

## {html_node}
## <a href="https://raleigh.craigslist.org/jwl/d/raleigh-vintage-hamilton-mens-automatic/77
## [1] <div class="title">Vintage Hamilton Men's Automatic Watch</div>
## [2] <div class="details">\n                <div class="price">$100</div>\ ...
```

Looks promising!

# Diving Deeper

```
# follow branching tree further:  
# title of listing  
title ← craigslist_listings ▷ html_elements("div.title") ▷ html_text2()  
  
# seems like 2 pieces of info stored in div.details  
details ← craigslist_listings ▷ html_elements("div.details")  
price ← details ▷ html_element("div.price") ▷ html_text2()  
location ← details ▷ html_element("div.location") ▷ html_text2()  
  
# we can use html_attr to grab the link to the listing  
link ← craigslist_listings ▷ html_attr("href")
```

# Taking a Look at Our Data

```
title[1]
```

```
## [1] "Vintage Hamilton Men's Automatic Watch"
```

```
price[1]
```

```
## [1] "$100"
```

```
location[1]
```

```
## [1] "Stonehenge, Raleigh NC"
```

```
link[1]
```

```
## [1] "https://raleigh.craigslist.org/jwl/d/raleigh-vintage-hamilton-mens-automatic/779607"
```

Looking *really* good



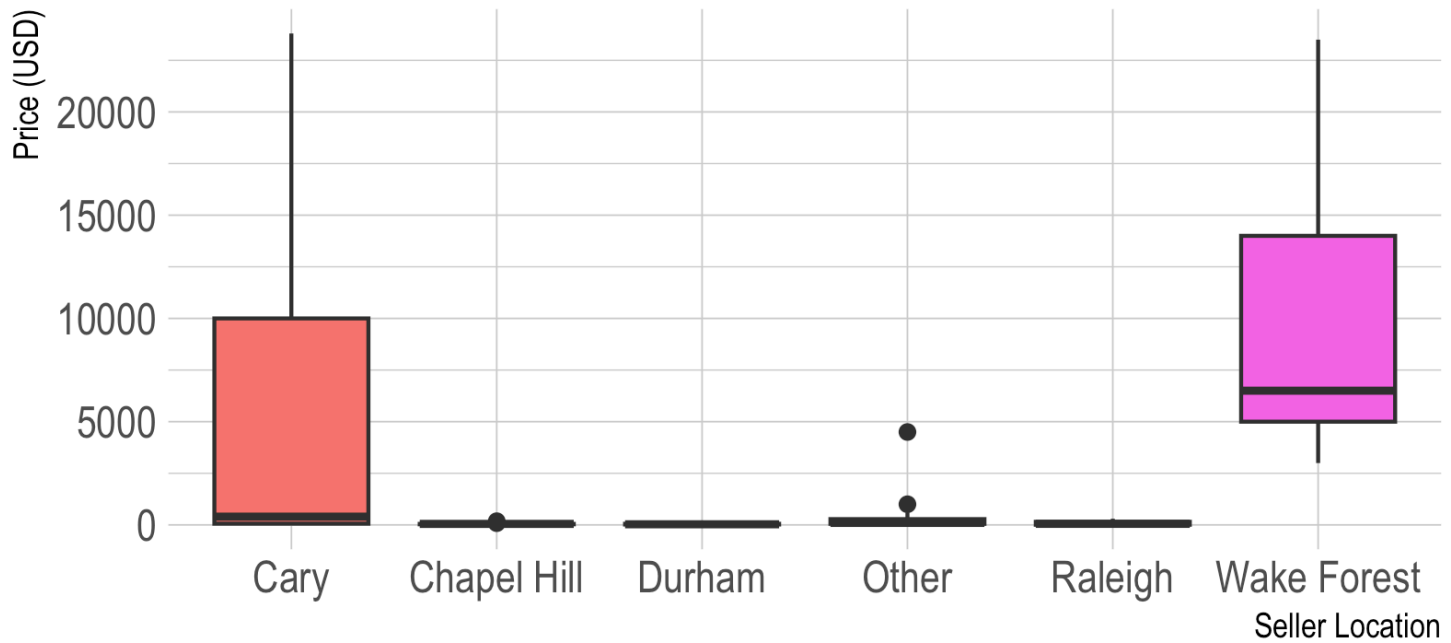
# Stick them in a data frame, clean up

```
watch_data <- data.frame(title,price,location,link) ▷  
  mutate(price = parse_number(price), # get rid of $s  
    clean_location = case_when(  
      grepl("durham",location, ignore.case=TRUE) ~ "Durham",  
      grepl("cary|apex",location, ignore.case=TRUE) ~ "Cary",  
      grepl("raleigh",location, ignore.case=TRUE) ~ "Raleigh",  
      grepl("chapel hill",location, ignore.case=TRUE) ~ "Chapel Hill",  
      grepl("wake forest",location, ignore.case=TRUE) ~ "Wake Forest",  
      .default = "Other"  
    )) ▷  
  filter(price>0)
```

How about another plot?

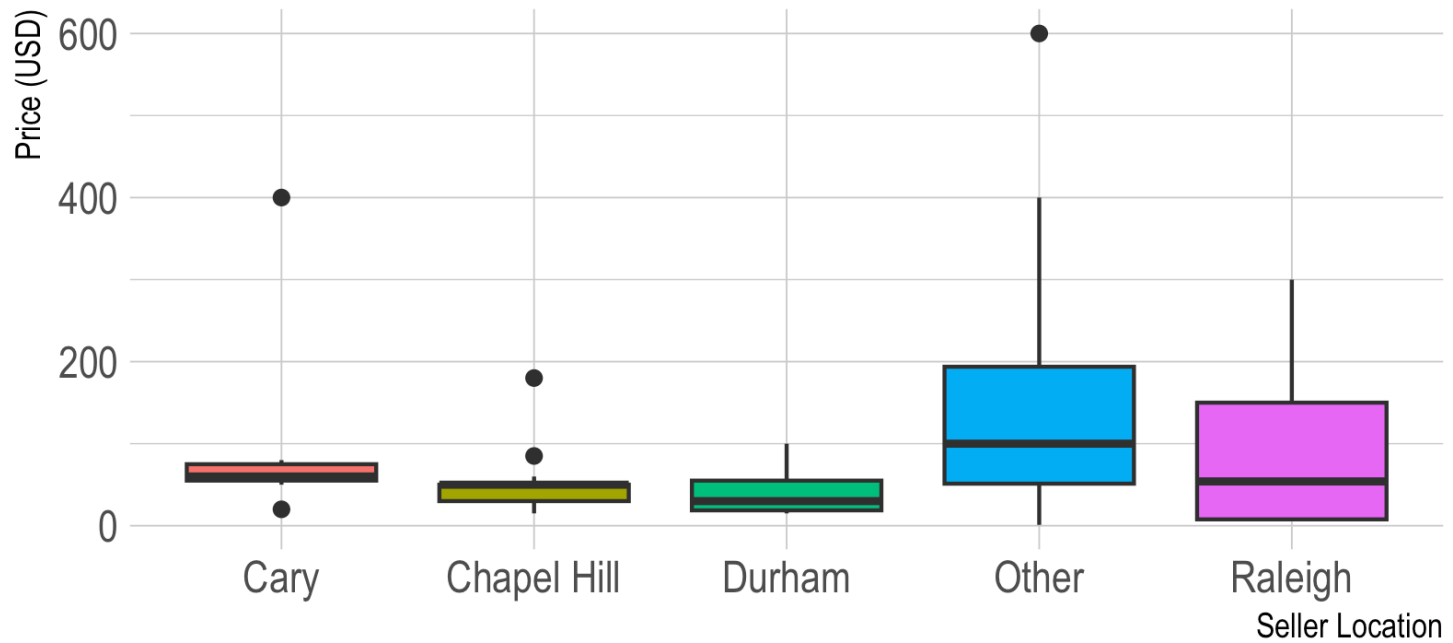
# Plot

## Watches for sale near Raleigh, NC



# Plot (for students)

## Watches for sale near Raleigh, NC



Sidenote: Wake Forest is gone!

# Summary

There's two ways to get data from the internet: server-side and client-side

We covered the server-side stuff today

To do so, we did some mucking around with CSS selectors

It is an art, not a science.

# Another Tool

Our approach worked pretty well for Craigslist, but we also got fairly lucky: there was only one page of search results when I wrote these notes.

It turns out you can use R to drive your actual computer, clicking on buttons (like Next Page) and so on.

You can do this using **RSelenium**, which is a huge pain but sometimes incredibly cool.

# Next Class

We'll be downloading macroeconomic data from FRED

Before class ***please*** make an account with FRED

<https://fredaccount.stlouisfed.org/login/secure/>

and obtain an API key <https://fredaccount.stlouisfed.org/apikey!!!>

PS: your api key is a secret, don't share it with people

# Next lecture: Scraping Client-Side with APIs

---