

Data Science for Economists


Lecture 13: Data Visualization - part 2, Maps

Drew Van Kuiken

University of North Carolina | ECON 370

Maps

Sample of New York City trees



Source: NYC Open Data

* Slides adapted from Grant McDermott's EC 607 at University of Oregon.

Maps in R

Packages for today:

```
## Load and install the packages that we'll be using today  
if (!require("pacman")) install.packages("pacman")  
pacman::p_load(sf, tidyverse, tigris, tidycensus, jsonlite, httr, maps, spData)
```

New packages: sf, tigris, tidycensus

Features

We'll use the `sf` and `tigris` packages to get our map data today. Be forewarned: this is a quick and dirty exposition, enough to get you and up and running, but not a full treatment of maps.

`sf`:

- Stands for *simple features*
- Treats features (e.g., zip codes, borders, trees, you name it) as data frames
- Great practice for data wrangling!

Shapefiles

The sf package has some maps pre-loaded for us. Let's start by looking at the counties of North Carolina.

```
file_loc = system.file("shape/nc.shp", package="sf")
nc = st_read(file_loc, quiet = TRUE)
head(nc)
```

```
## Simple feature collection with 6 features and 14 fields
```

```
## Geometry type: MULTIPOLYGON
```

```
## Dimension:      XY
```

```
## Bounding box:  xmin: -81.74107 ymin: 36.07282 xmax: -75.77316 ymax: 36.58965
```

```
## Geodetic CRS:  NAD27
```

```
##      AREA PERIMETER CNTY_ CNTY_ID      NAME  FIPS FIPSNO CRESS_ID BIR74 SID74
## 1  0.114      1.442  1825   1825      Ashe 37009  37009         5  1091    1
## 2  0.061      1.231  1827   1827 Alleghany 37005  37005         3   487    0
## 3  0.143      1.630  1828   1828      Surry 37171  37171        86  3188    5
## 4  0.070      2.968  1831   1831 Currituck 37053  37053        27   508    1
## 5  0.153      2.206  1832   1832 Northampton 37131  37131        66  1421    9
## 6  0.097      1.670  1833   1833  Hertford 37091  37091        46  1452    7
##      NWBIR74 BIR79 SID79 NWBIR79      geometry
```

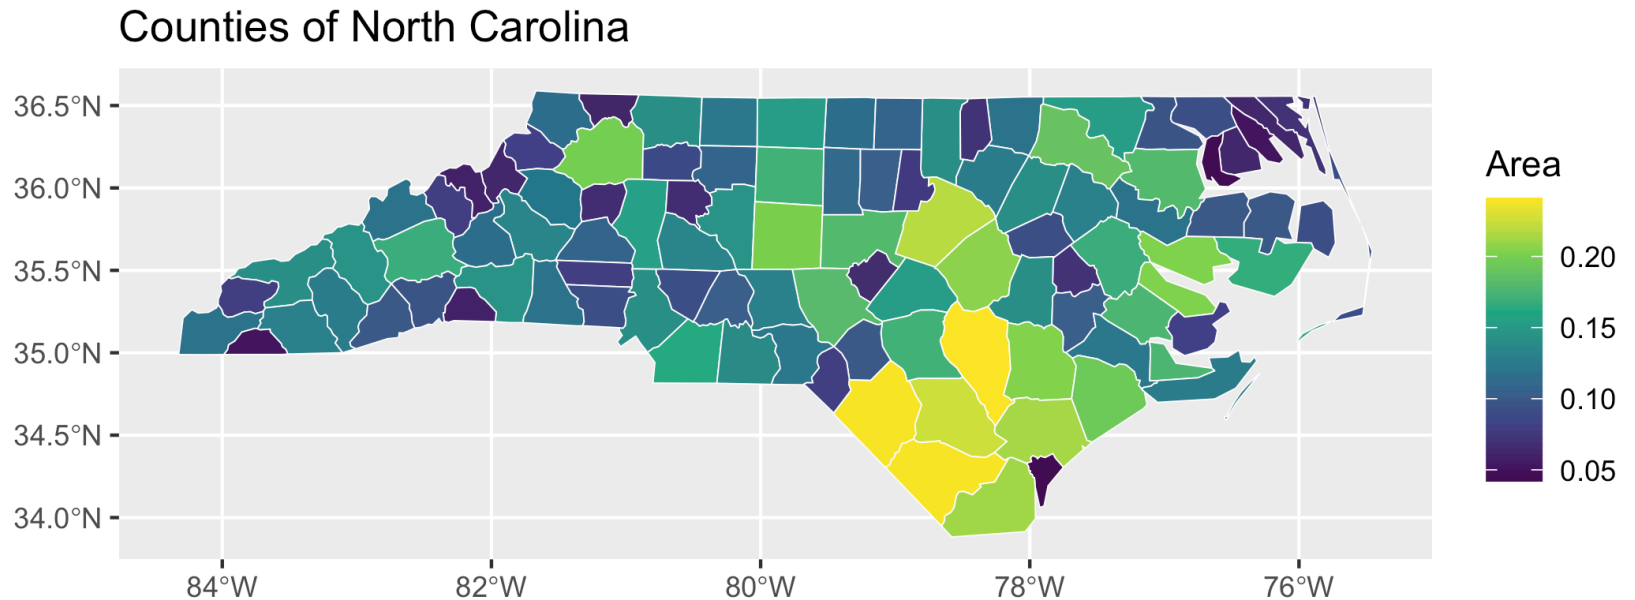
```
## 1      10   1364     0      19 MULTIPOLYGON (((-81.47276 3 ...
```

```
## 2      10    542     3      12 MULTIPOLYGON (((-81.23989 3 ...
```

```
## 3     208   3616     6     260 MULTIPOLYGON (((-80.45634 3 ...
```

Our first plot

```
ggplot(nc) + geom_sf(aes(geometry=geometry, fill=AREA), col="white") + scale_fill_vir:
```



Geodetic CRS

Otherwise known as the **projection** we're using. There are **lots** of projections we can use.

To change projection, we can use the `st_transform` function:

```
nc_moll <- nc ▷  
  st_transform(crs = "+proj=moll") # using the Mollweide projection  
head(nc_moll)
```

```
## Simple feature collection with 6 features and 14 fields
```

```
## Geometry type: MULTIPOLYGON
```

```
## Dimension:      XY
```

```
## Bounding box:  xmin: -7160488 ymin: 4345895 xmax: -6638106 ymax: 4404766
```

```
## Projected CRS: +proj=moll
```

```
##      AREA PERIMETER CNTY_ CNTY_ID      NAME  FIPS FIPSNO CRESS_ID BIR74 SID74  
## 1 0.114      1.442  1825   1825      Ashe 37009  37009         5  1091    1  
## 2 0.061      1.231  1827   1827 Alleghany 37005  37005         3   487    0  
## 3 0.143      1.630  1828   1828      Surry 37171  37171        86  3188    5  
## 4 0.070      2.968  1831   1831 Currituck 37053  37053        27   508    1  
## 5 0.153      2.206  1832   1832 Northampton 37131  37131        66  1421    9  
## 6 0.097      1.670  1833   1833   Hertford 37091  37091        46  1452    7
```

```
##      NWBIR74 BIR79 SID79 NWBIR79      geometry
```

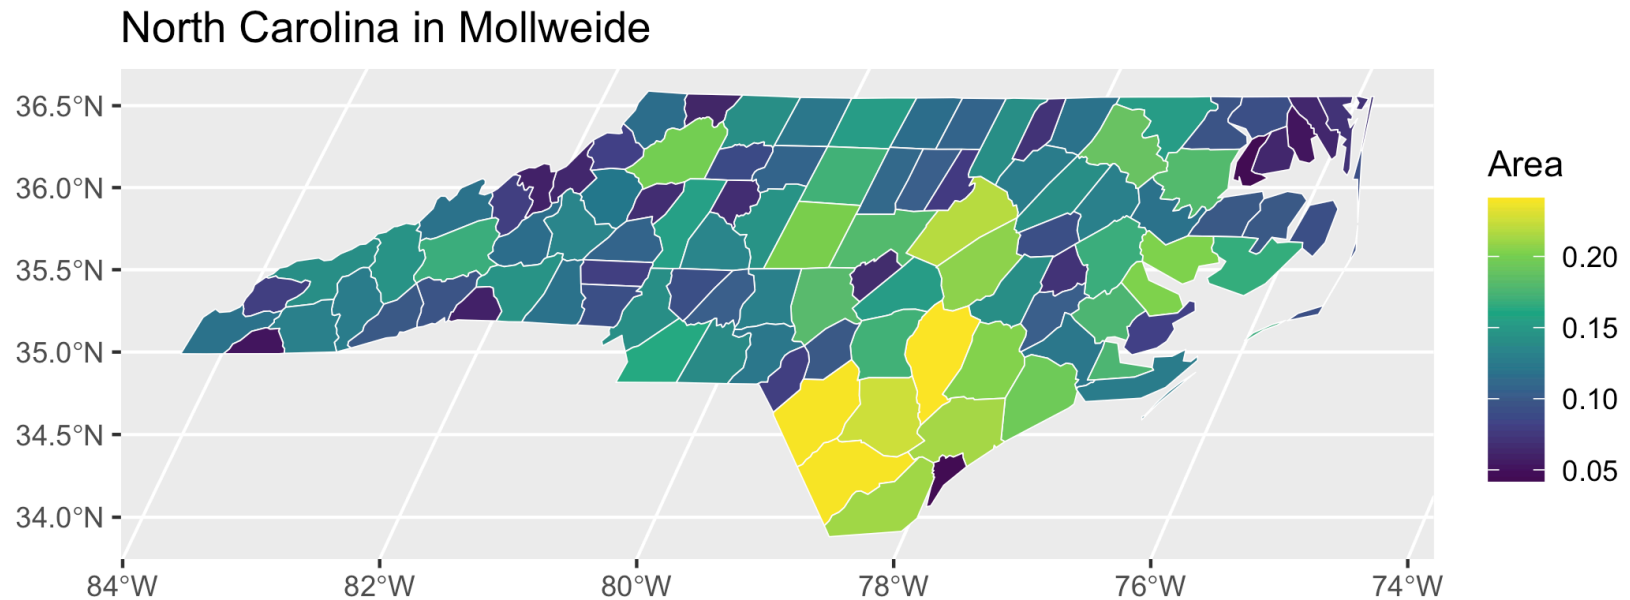
```
## 1         10  1364     0        19 MULTIPOLYGON (((-7145982 43 ...
```

```
## 2         10   542     3        12 MULTIPOLYGON (((-7118092 43 ...
```

```
## 3         208  2616     6        260 MULTIPOLYGON ((( 7056268 42 ...
```

Mollweide?

Looks like this:



Not super useful. Keep the projection thing in mind though, it will come up later.

Wrangling

We can use the tidyverse to work with sf objects:

```
triangle <- nc ▷  
  filter(NAME %in% c("Durham", "Wake", "Orange", "Chatham")) ▷  
  mutate(AREA = AREA*1000) ▷  
  select(NAME, AREA, everything())  
head(triangle)
```

```
## Simple feature collection with 4 features and 14 fields
```

```
## Geometry type: MULTIPOLYGON
```

```
## Dimension:      XY
```

```
## Bounding box:   xmin: -79.55536 ymin: 35.51024 xmax: -78.25455 ymax: 36.23569
```

```
## Geodetic CRS:  NAD27
```

```
##      NAME AREA PERIMETER CNTY_ CNTY_ID  FIPS FIPSNO CRESS_ID BIR74 SID74
```

```
## 1  Orange  104      1.294  1907   1907 37135  37135      68  3164    4
```

```
## 2  Durham   77      1.271  1908   1908 37063  37063      32  7970   16
```

```
## 3   Wake   219      2.130  1938   1938 37183  37183      92 14484   16
```

```
## 4 Chatham  180      2.142  1973   1973 37037  37037      19  1646    2
```

```
##  NWBIR74 BIR79 SID79 NWBIR79 geometry
```

```
## 1      776  4478     6    1086 MULTIPOLYGON (((-79.01814 3 ...
```

```
## 2     3732 10432    22    4948 MULTIPOLYGON (((-79.01814 3 ...
```


```
## 3     4397 20857    31    6221 MULTIPOLYGON (((-78.92107 3 ...
```

```
## 4      591  2398     3     687 MULTIPOLYGON (((-79.55536 3 ...
```

Special Operations

We can melt geographies into larger geographic units using `st_union()`:

```
triangle ▷ st_union() ▷ ggplot() + geom_sf(fill=NA,col="black") + labs(title = "The
```



Other Operations

We're not going to focus on the full suite of operations available today. Here's a quick overview:

- `st_area`: get area of shape
- `st_centroid`: get geographic center of area
- `st_boundary`: returns boundary of geometry
- `st_buffer`: creates buffer around each observation
- `st_distance`: distance between two objects (or more)

Multiple Datasets

Let's pick up two new datasets:

```
france = st_as_sf(map('france', plot=FALSE, fill=TRUE))
france = france[-c(95),] # observation causing some issues, not important
head(france)
```

```
## Simple feature collection with 6 features and 1 field
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: 0.06215676 ymin: 48.86568 xmax: 5.372333 ymax: 51.09752
## Geodetic CRS:   +proj=longlat +ellps=clrk66 +no_defs +type=crs
##               ID                                geom
## Nord              Nord MULTIPOLYGON (((2.557093 51 ...
## Pas-de-Calais     Pas-de-Calais MULTIPOLYGON (((2.105322 51 ...
## Somme              Somme MULTIPOLYGON (((1.623103 50 ...
## Ardennes          Ardennes MULTIPOLYGON (((4.220728 49 ...
## Seine-Maritime    Seine-Maritime MULTIPOLYGON (((1.419646 50 ...
## Aisne              Aisne MULTIPOLYGON (((3.15867 50....
```

```
data("seine", package = "spData")
head(seine)
```

```
## Simple feature collection with 3 features and 1 field
```

Merging

A couple thoughts:

1. `st_as_sf`: turns non-shapefile data into shapefile data.
`st_as_sf(x, coordinates=c("longitude", "latitude"))` turns any lat/lon data into a sf object.
2. `st_intersect`: creates a dataset with the exact points of overlap between two objects.
3. `st_join`: joins together two sf objects based on any overlapping geometries
 - Example of 2 and 3 in a sec

Try taking the intersection of the france and seine sf objects now.

Didn't Work, Huh


Why didn't it work?

```
seine_crs = st_transform(seine, crs = st_crs(france))  
france_intersected = st_intersection(france, seine_crs)
```

```
## Warning: attribute variables are assumed to be spatially constant throughout  
## all geometries
```

The projections were different!

So what'd we create?




Trying out joins

```
france_river <- st_join(france,seine_crs) ▷ # what kind of join is this?  
  filter(!is.na(name)) ▷  
  distinct(ID, .keep_all = T) # some rows merge twice because 2 branches of river  
head(france_river)
```

```
## Simple feature collection with 6 features and 2 fields  
## Geometry type: MULTIPOLYGON  
## Dimension:      XY  
## Bounding box:   xmin: 0.06215676 ymin: 48.44897 xmax: 5.032294 ymax: 50.08539  
## Geodetic CRS:   +proj=longlat +ellps=clrk66 +no_defs +type=crs  
##               ID  name                                           geom  
## Seine-Maritime Seine-Maritime Seine MULTIPOLYGON (((1.419646 50 ...  
## Aisne           Aisne Marne MULTIPOLYGON (((3.15867 50....  
## Eure           Eure Seine MULTIPOLYGON (((0.4339198 4 ...  
## Marne          Marne Marne MULTIPOLYGON (((4.059515 49 ...  
## Val-Doise      Val-Doise Seine MULTIPOLYGON (((1.706263 49 ...  
## Yvelines       Yvelines Seine MULTIPOLYGON (((1.624106 49 ...
```


Put It In a Graph



tidycensus

The tidycensus package lets you avoid setting up even an API call to get census data, it's sort of amazing.

We'll only scratch the surface of the tidycensus package. If you wanted to get ACS data (zipcode-level survey on all kinds of stuff), you'd need to grab an API key from the census. The instructions to do this are in the class script for today.

One thing that's available in the tidycensus: TIGER/Line shapefiles (to access these, we loaded the `tigris` package)

Tiger shapefiles can give us shapefiles for things like: nations, regions, states, counties, census tracts, school districts, zip codes, and **way more**

Zip Codes in NY

Let's grab the zip codes for the urban area New York City:

```
ny_zips <- zctas(state = "NY", class = "sf", year = 2010) # zips only available in 20.  
urb <- urban_areas(year=2020) > filter(grepl("New York",NAME10))  
ny_urb_zips <- st_join(ny_zips,urb) > filter(!is.na(NAME10))
```

Your task -- create 2 graphs:

1. Zip codes of NYC which contain trees downloaded from NYC Open Data
2. All zip codes in NY urban area, with points representing trees in NYC